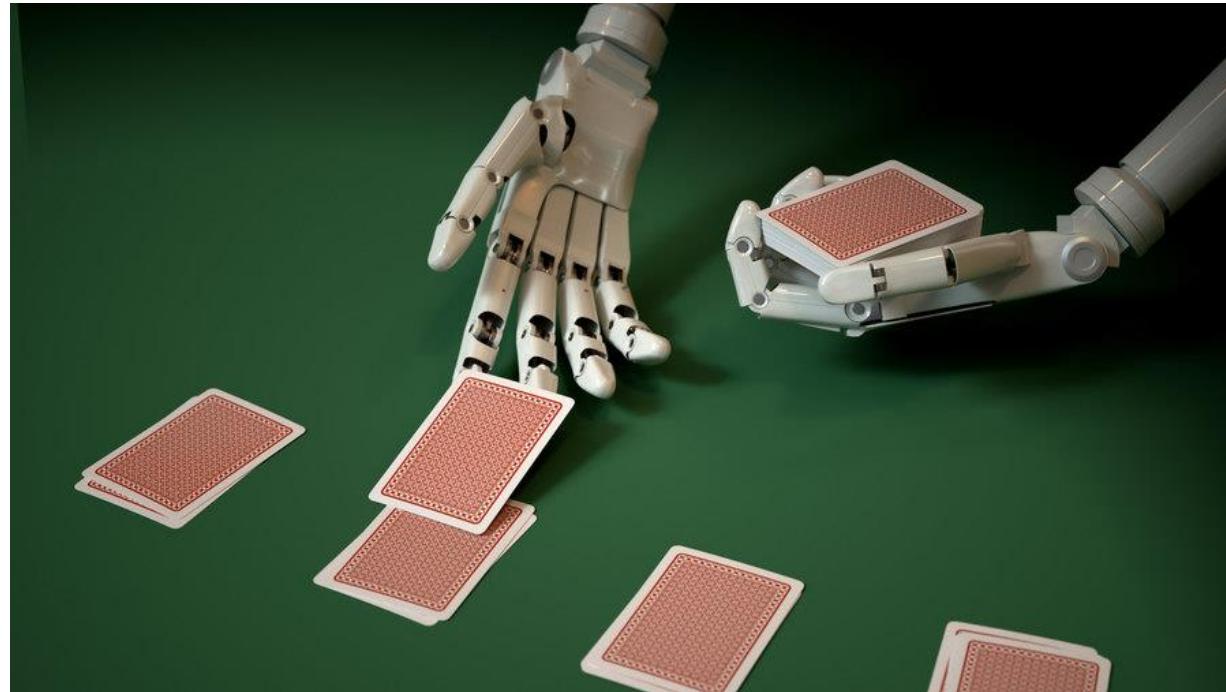


practical deep learning introduction to machine learning

Alex Honchar
University of Verona

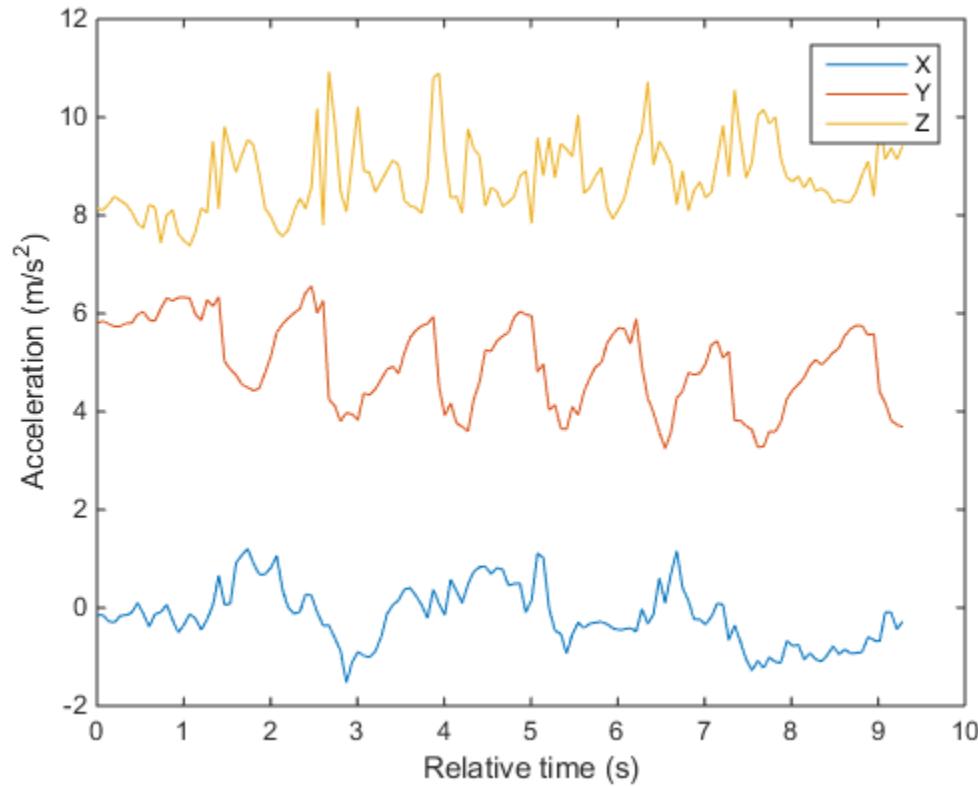
my experience

Playing bot for online game (UA, RU 2014)



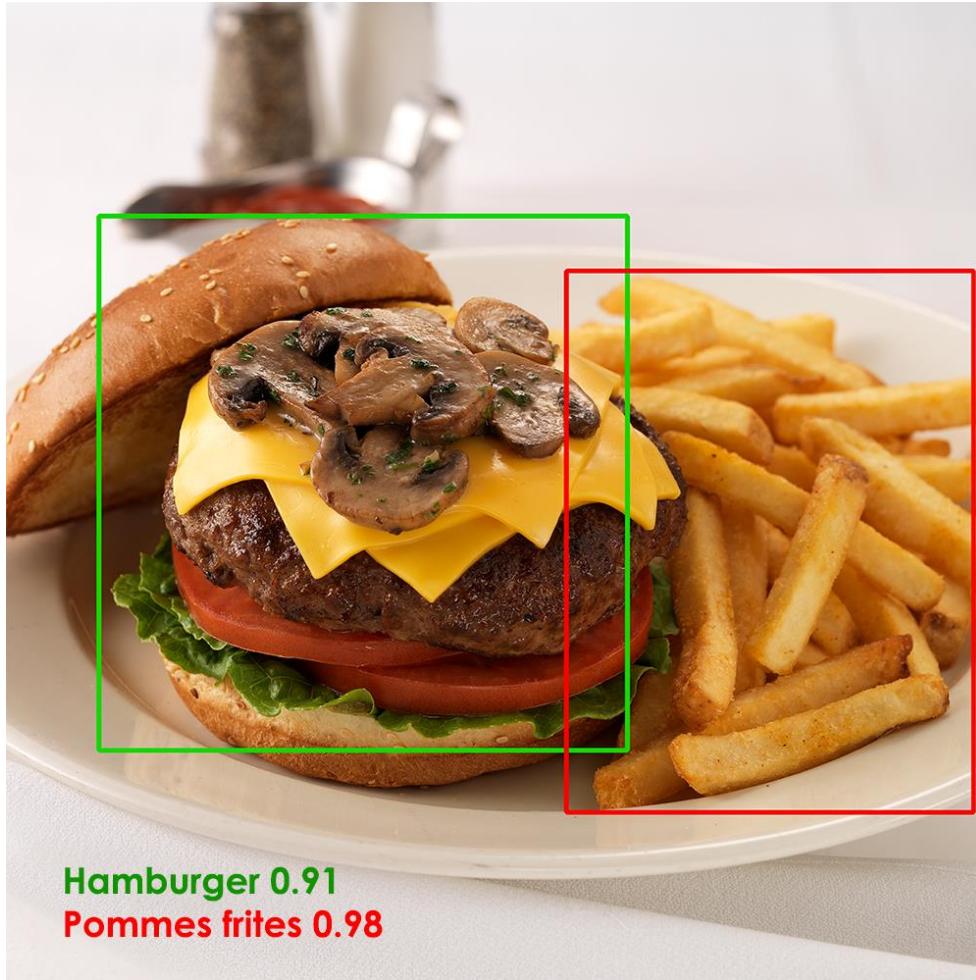
- Game modelling
- Opponent action prediction

Mobile sensor data analysis (USA 2015)



- Activity recognition
- Blood pulse calculation

Food image recognition (USA 2015)



- Recognition of several meals
- Recognition of ingredients inside

Advertisement classification (RU 2016)

Selling used smartphone
Meizu M1 Note for \$100



Selling used **smartphone**
Meizu **M1 Note** for \$100

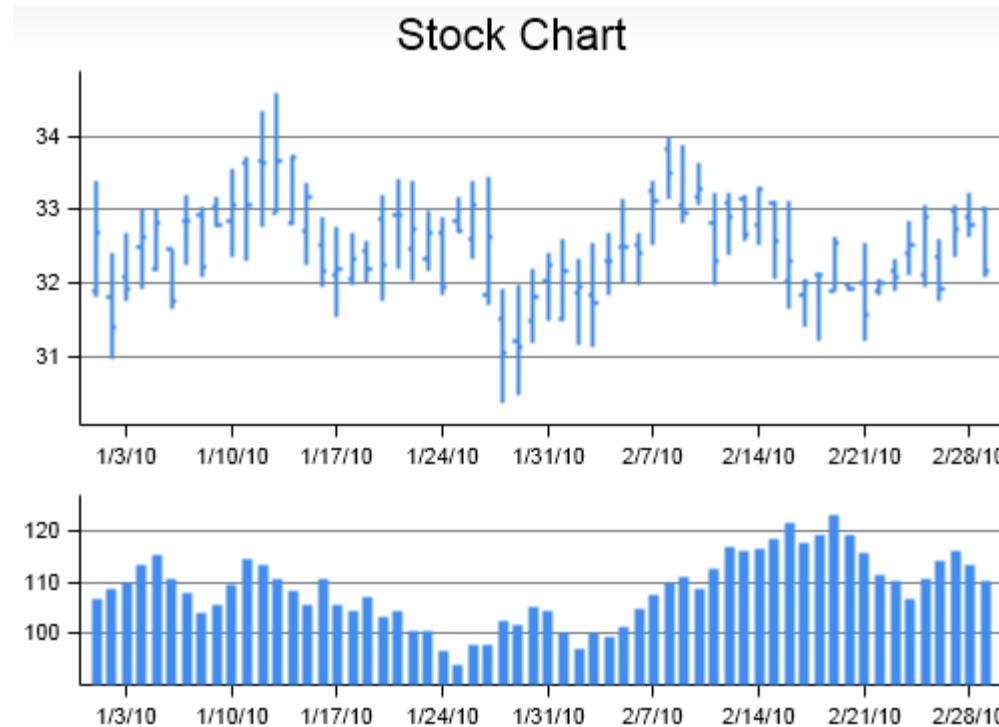
- Named entity recognition (NER) problem
- Over 100'000'000 advertisements

Black - not defined entity
Green - brand
Red - category
Blue - model

Image style transfer (mlvch.com 2016)

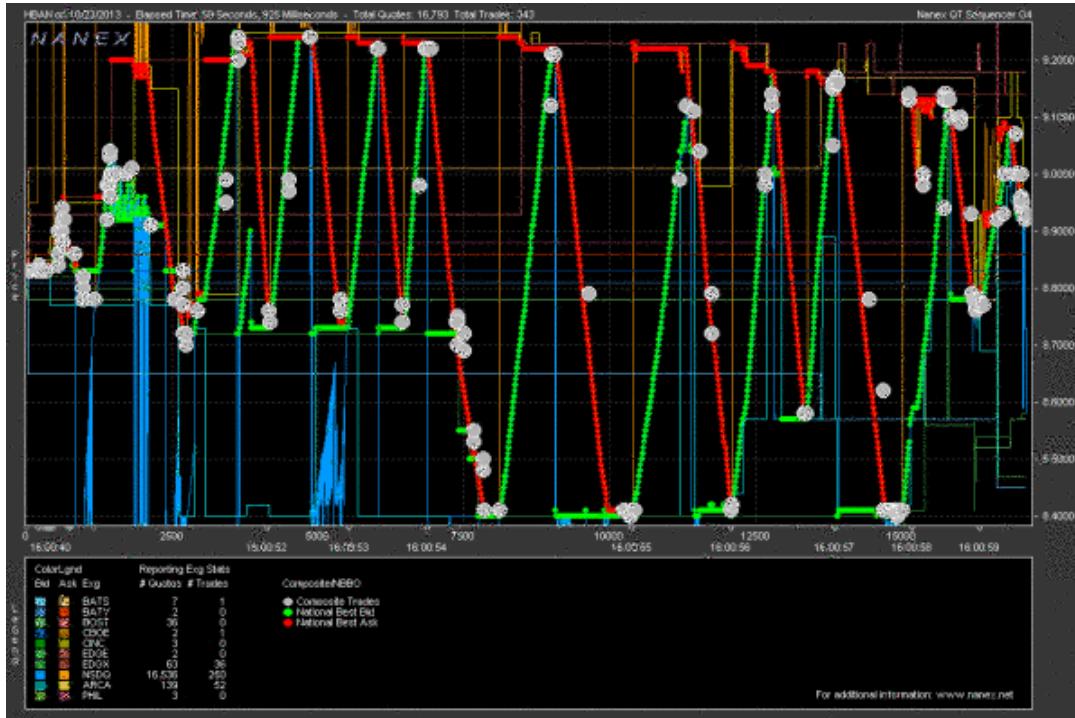


Financial forecasting challenge (2016)



- Numer.ai – forecasting from encrypted financial data
- In September achieved top 10% in public leaderboard

Current research



High frequency trading



Bots
for Messenger

Conversational bots

my mission

machines **work**, people **create**

my mission

machines **work**, people **create**
machines **learn**, people **discover**

my mission

machines **work**, people **create**
machines **learn**, people **discover**
machines **like**, people **love**

your experience

seminar goals

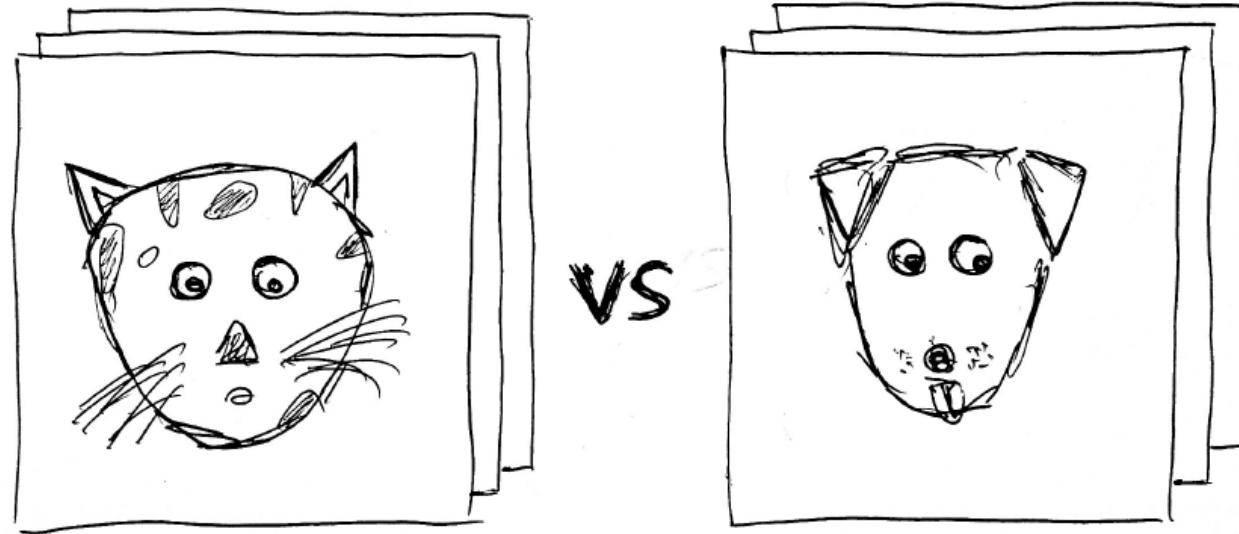
1. You understand how deep learning works
2. You create a prototype with hands-on tools
3. You know the path to get the job in machine learning

Day 1 goals

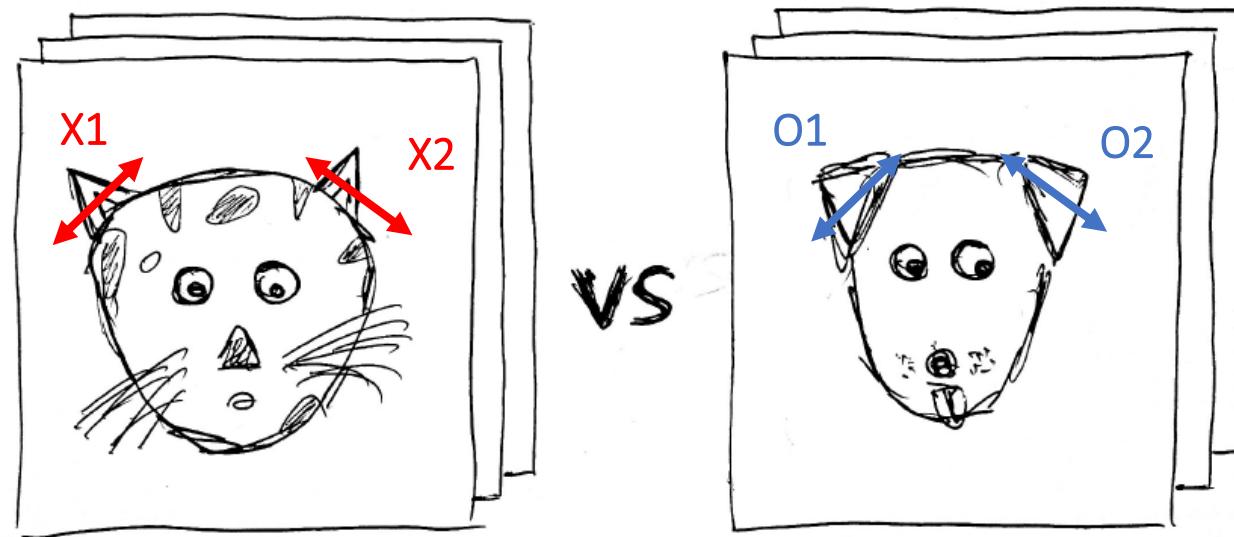
- You **understand** basics of machine learning
- You **understand** difference between machine learning and deep learning (most engineers don't!)
- You **can** define machine learning problem with ERTP blocks

machine learning in nutshell

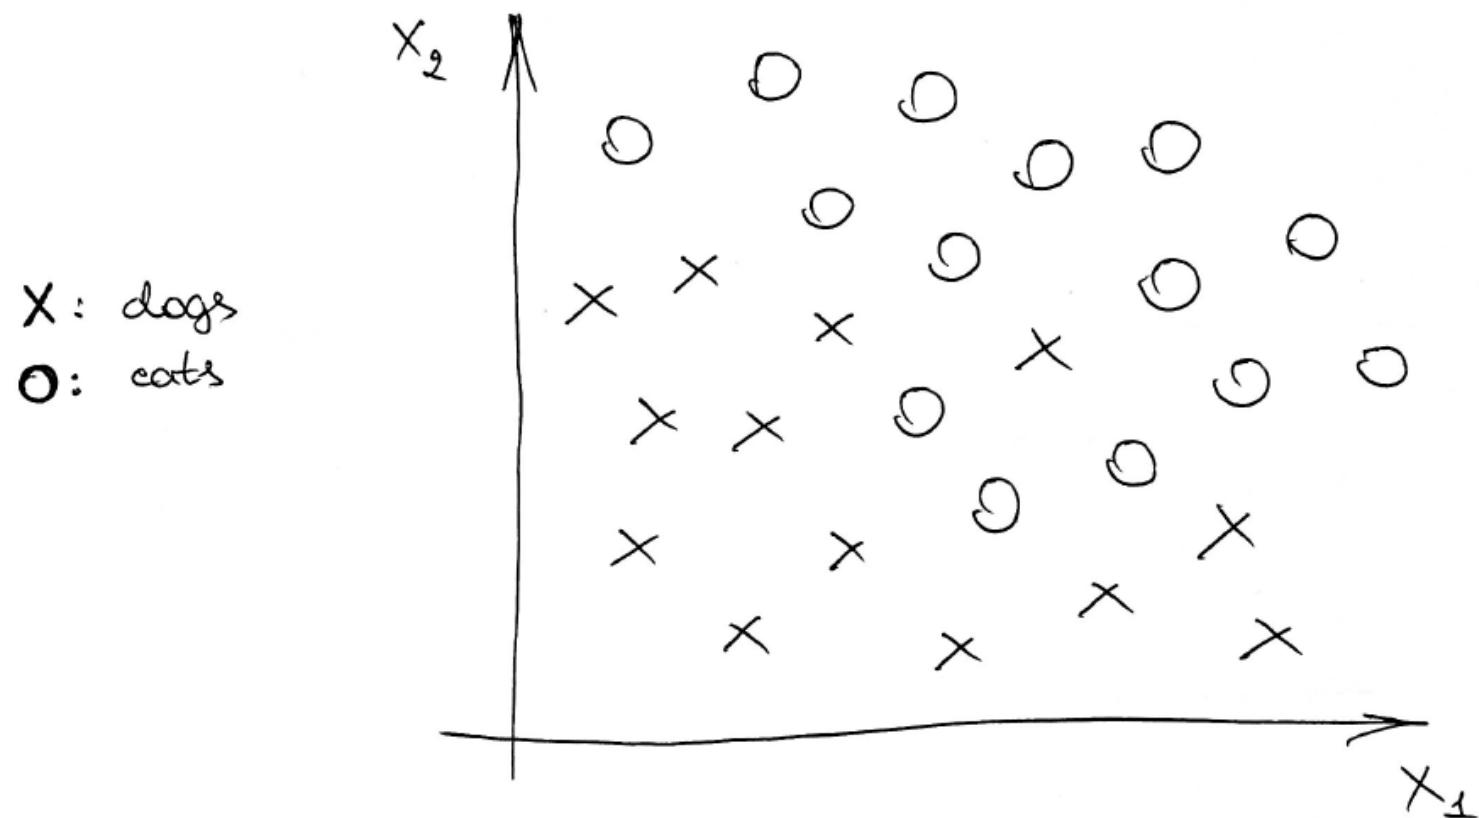
Cats vs Dogs



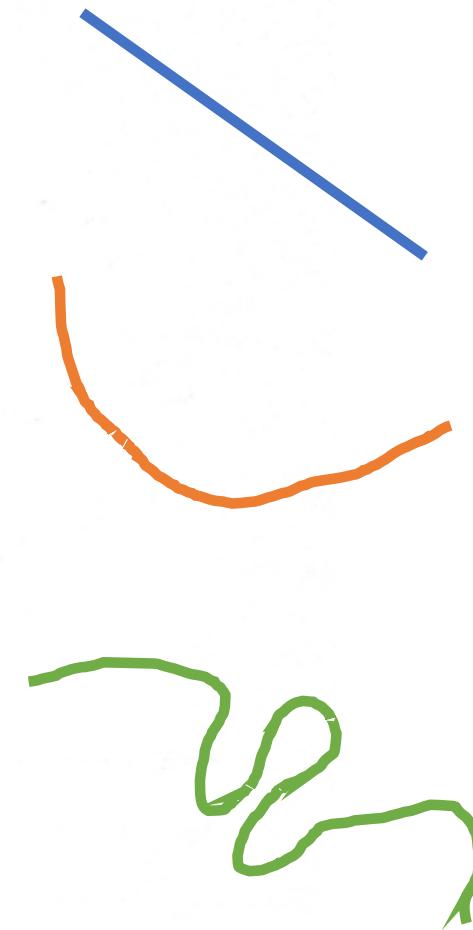
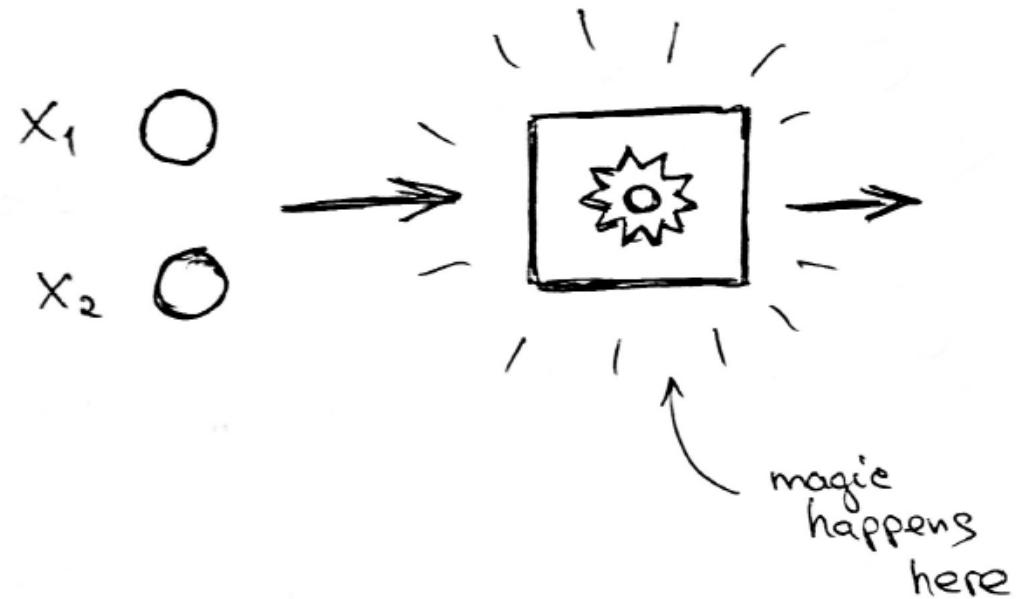
Cats vs Dogs



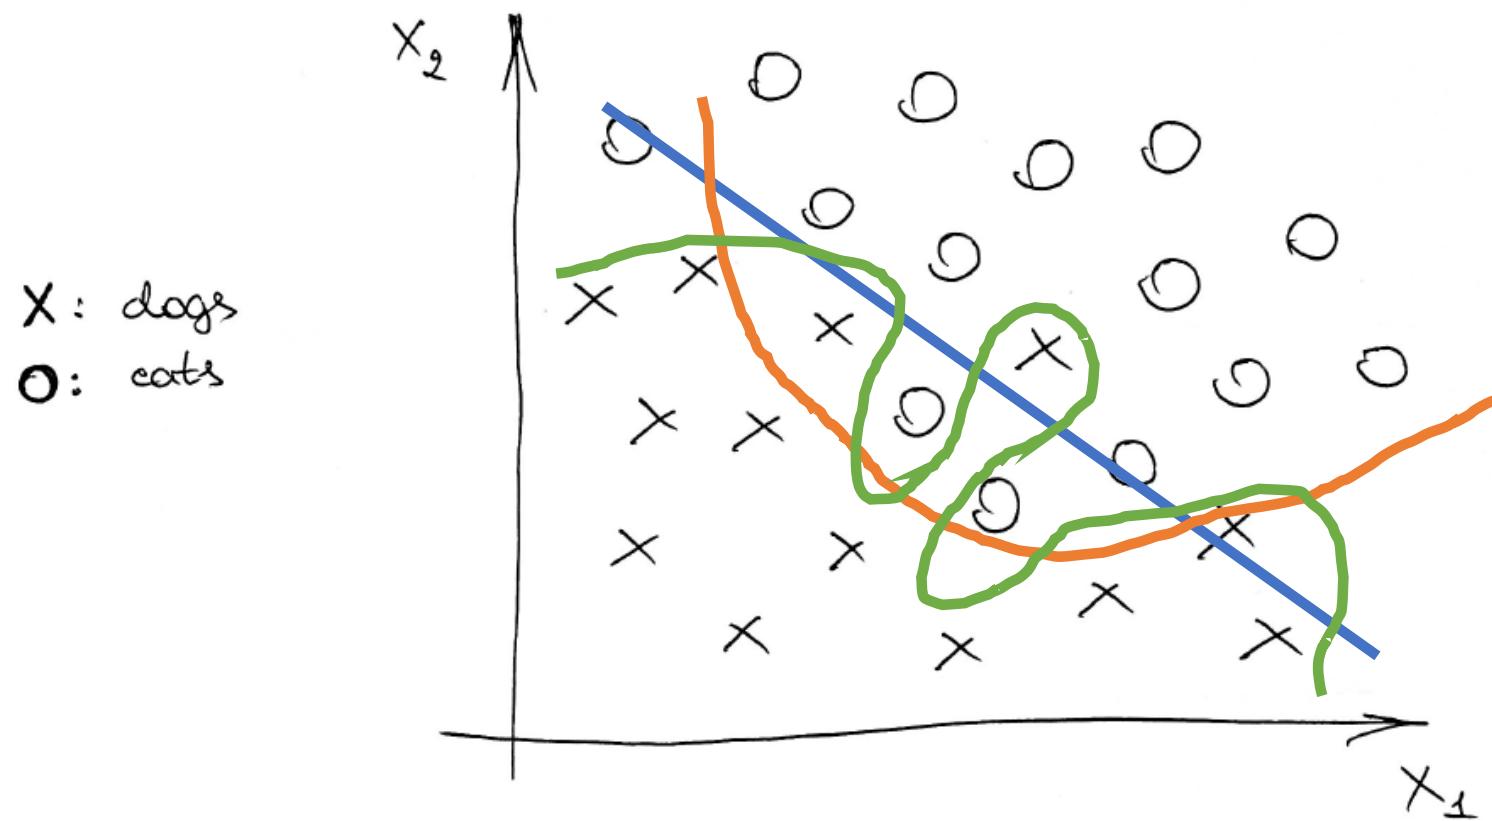
Visual machine learning



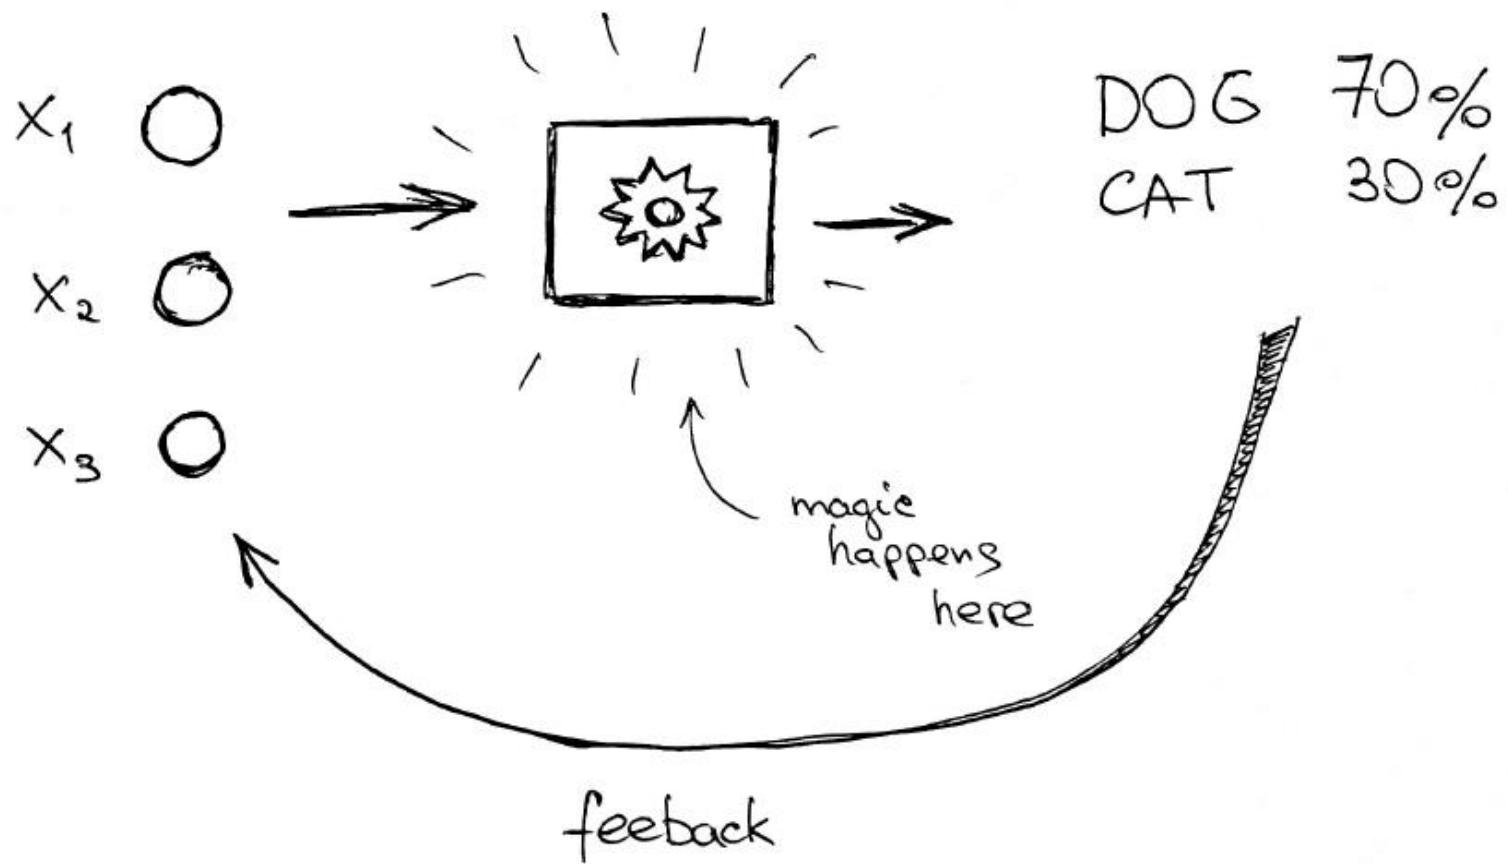
Visual machine learning



Cats vs Dogs



How it learns?



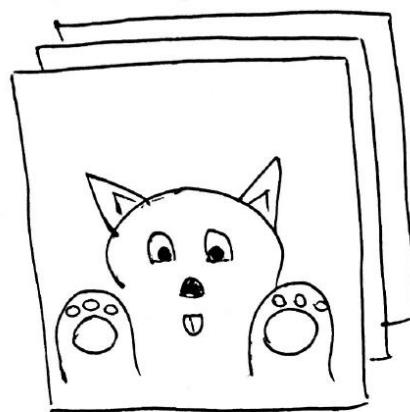
résumé

Learning algorithm is just a **baby** – we show him examples of different objects and he tries to **generalize** them to correctly recognize new objects (separate one from another)

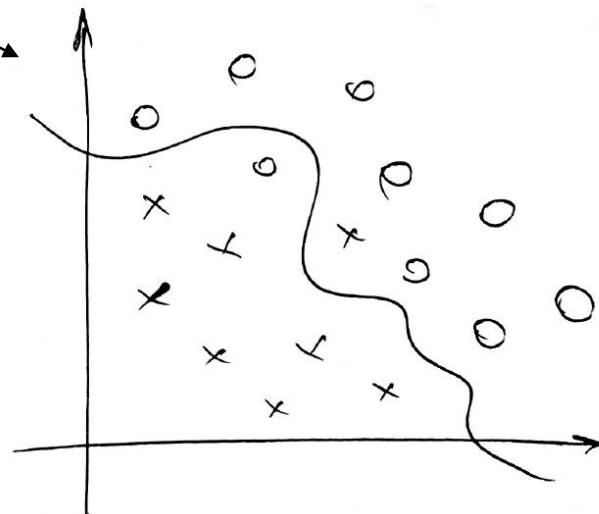
Mathematical framework

- Program is said to learn from **experience E** with respect to some class of tasks **T** and performance measure **P** if performance improves

Mitchell, 1997



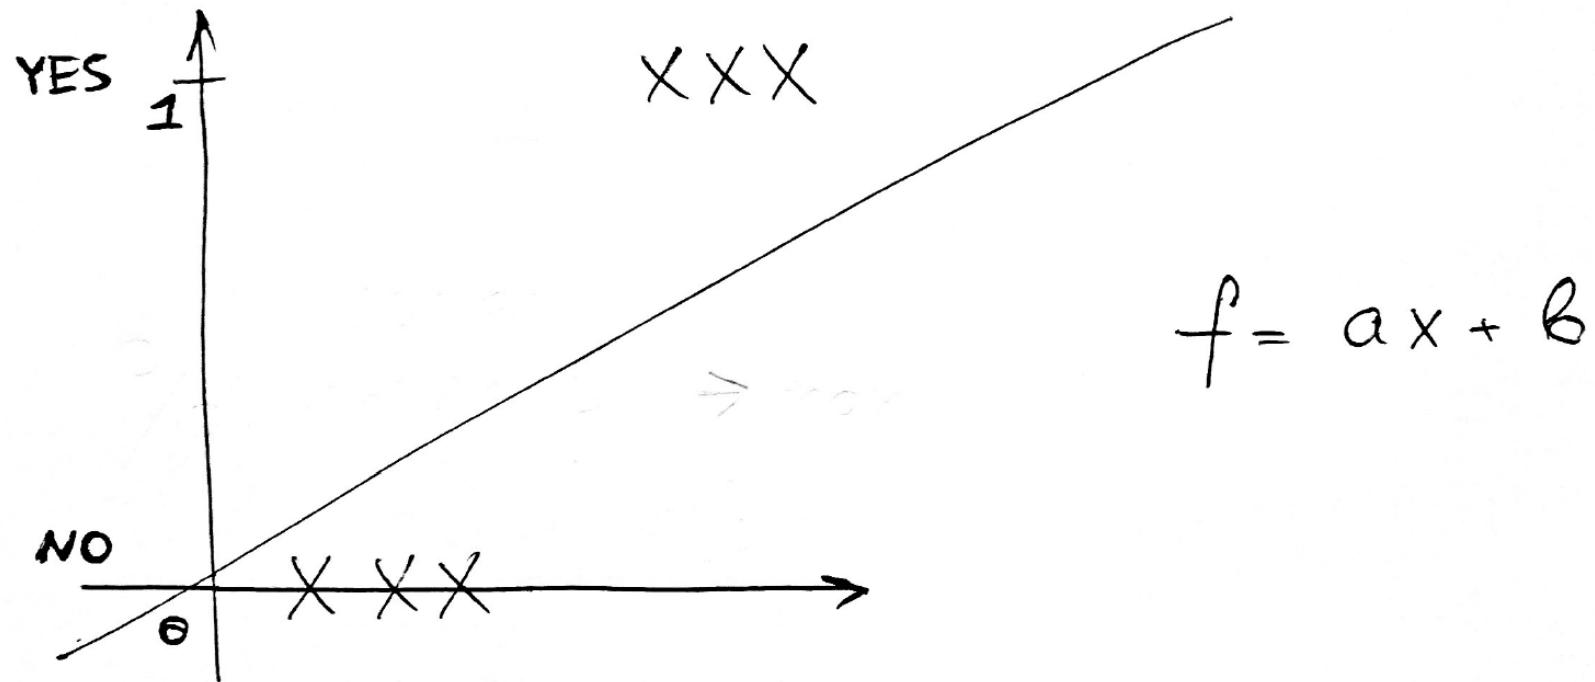
$$\{ R^{250 \times 250} \}$$



$$f: R^{250 \times 250} \rightarrow \{ \text{cat, dog} \}$$

% correct answers $\{ \text{cat, dog} \} \rightarrow \max$

Logistic regression



Logistic regression

Hypothesis

$$h_w(x) = W^T x$$

$$W_0 + W_1 x$$

$$W_0 + W_1 x_1 + W_2 x_2$$

$$W_0 + W_1 x_1 + W_2 x_2^2 + W_3 x_2^2$$

$$W_0 + W_1 x_1 + W_2 x_2 + W_3 x_2^2 + W_4 x_2^2$$

Threshold classifier

if $h_w(x) \geq 0.5$:

" $y = \text{YES}$ "

if $h_w(x) < 0.5$:

" $y = \text{NO}$ "

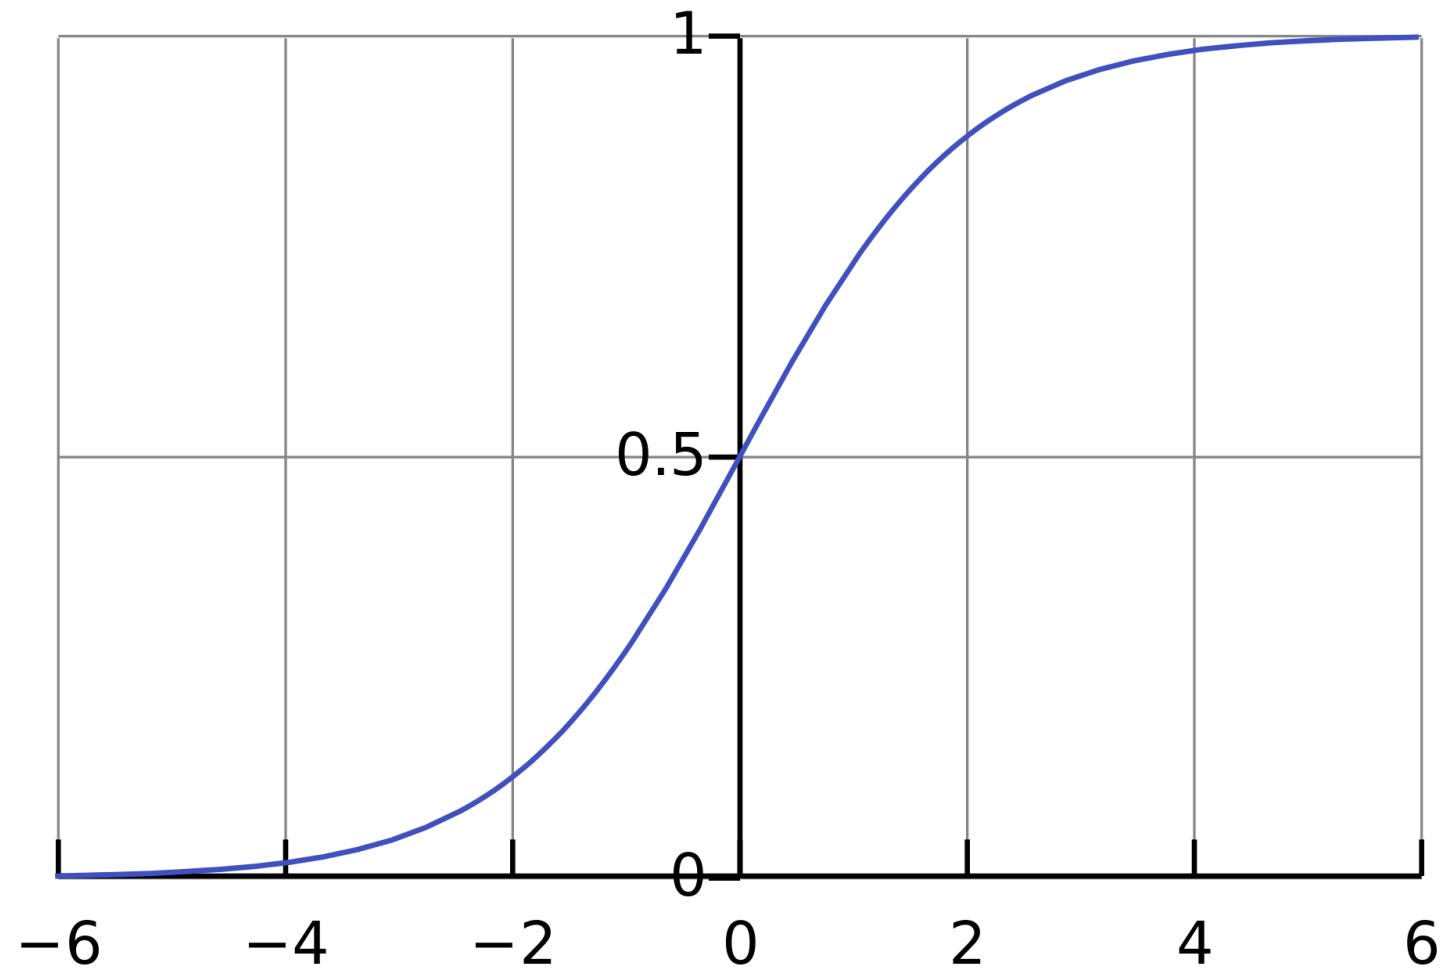
Logistic regression

$$h_w(x) \in [0, 1]$$

$$h_w(x) = g(W^T x)$$

$$g(z) = \frac{1}{e^{-z} + 1} \Rightarrow h_w(x) = \frac{1}{1 + e^{-W^T x}}$$

Logistic function



Logistic regression

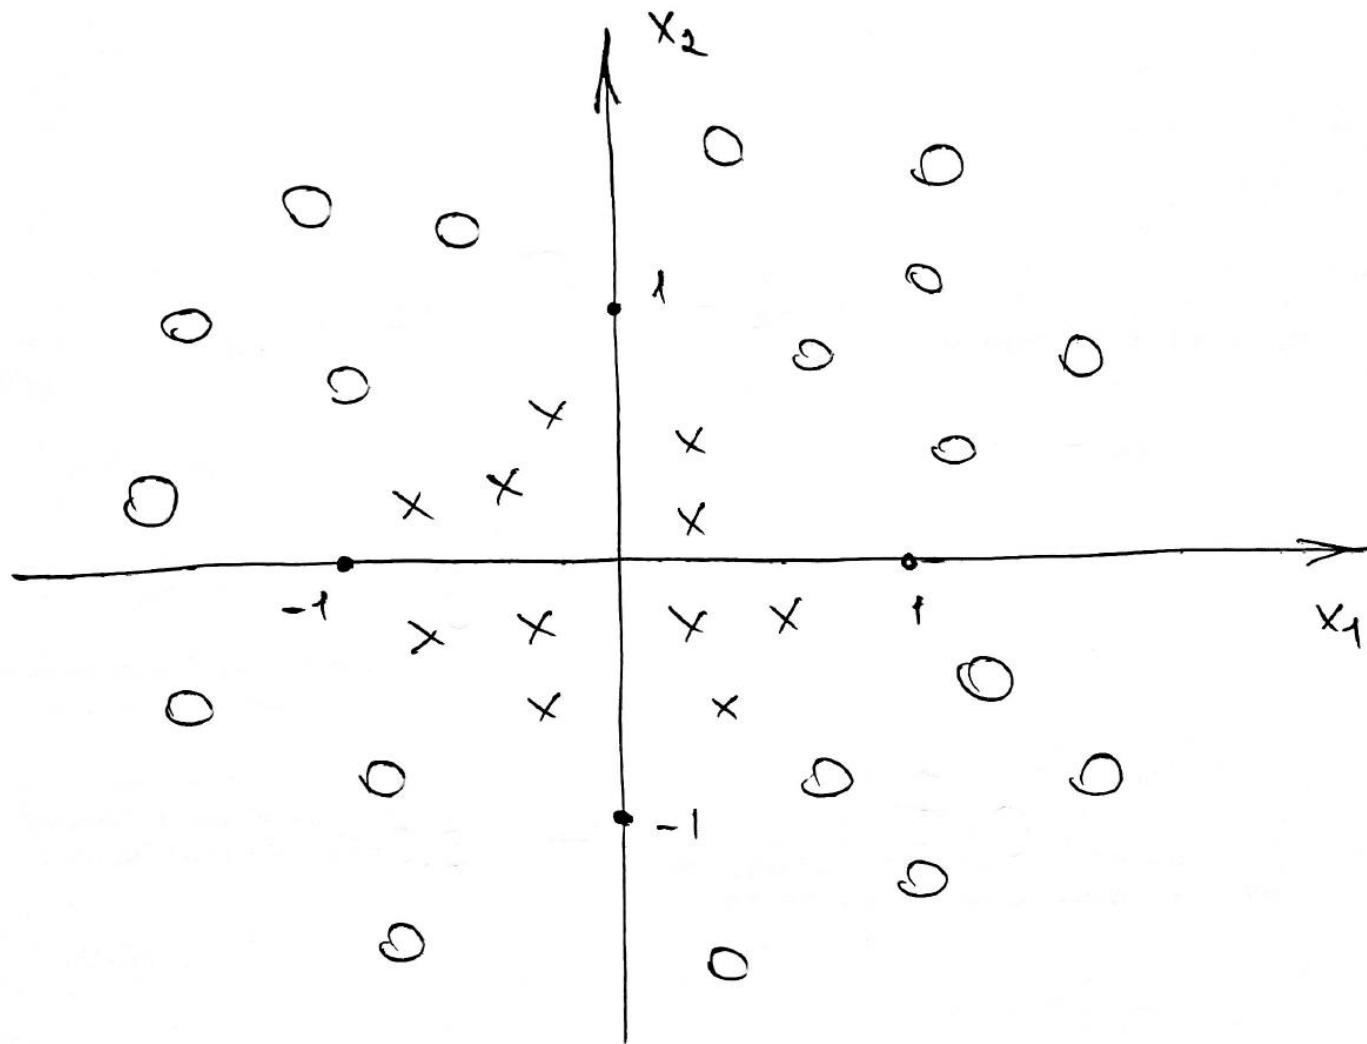
$$h_w(x) = P(y=1 | x; W)$$

\curvearrowright \curvearrowleft

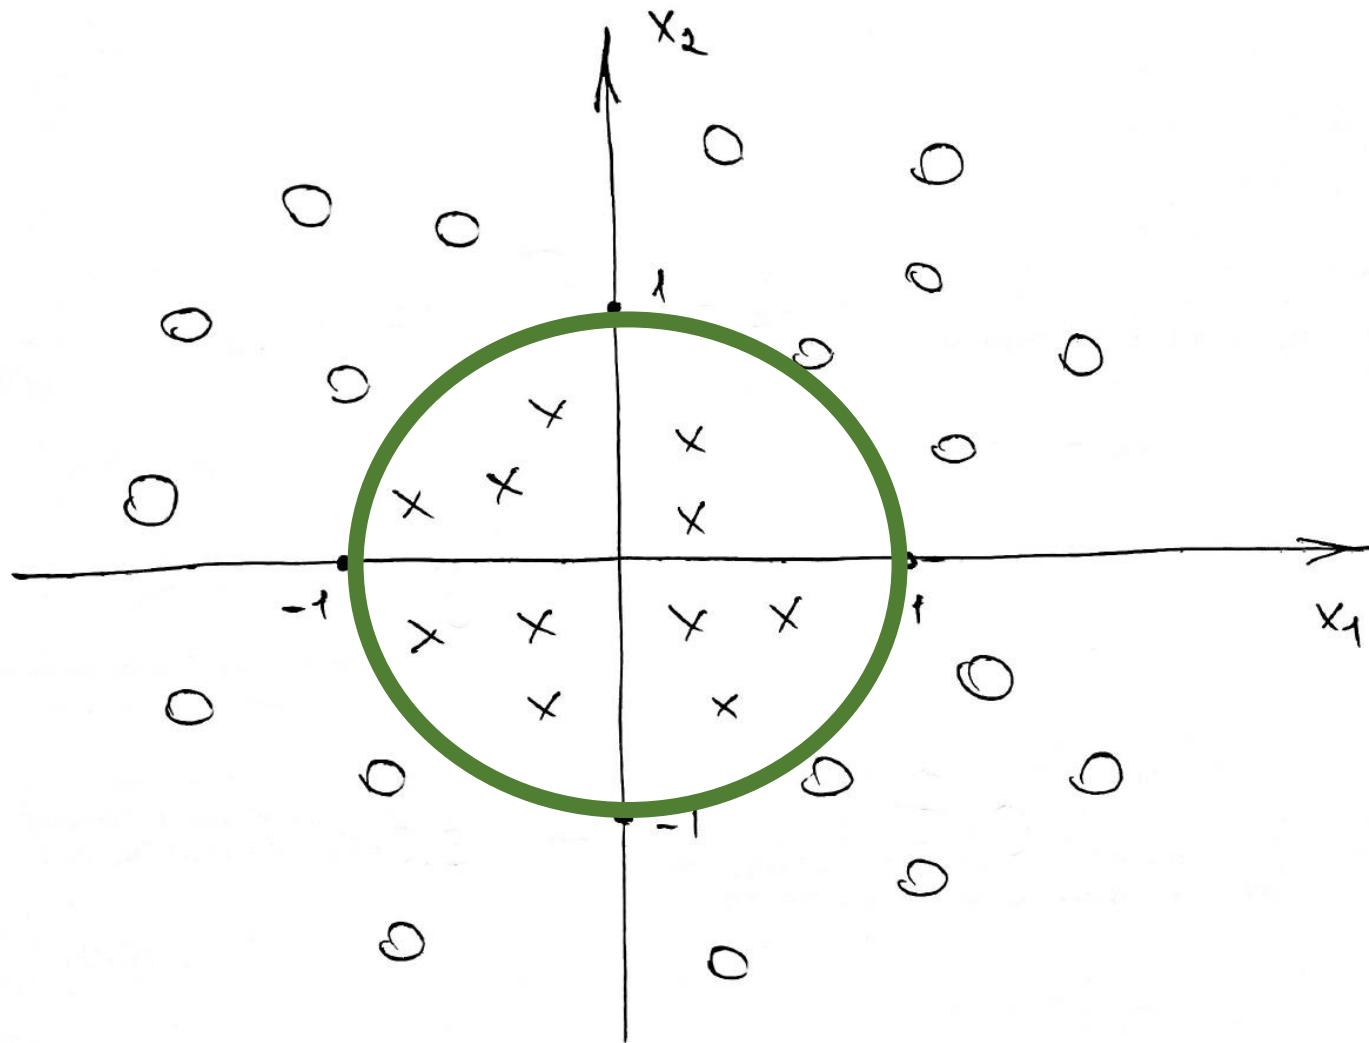
$$g(W^T x)$$

$$h_w(x) \geq 0.5 \Rightarrow W^T x \geq 0$$

Logistic regression



Logistic regression



Logistic regression

$$h_w(x) = g(W_0 + W_1x_1 + W_2x_2 + W_3x_1^2 + W_4x_2^2)$$

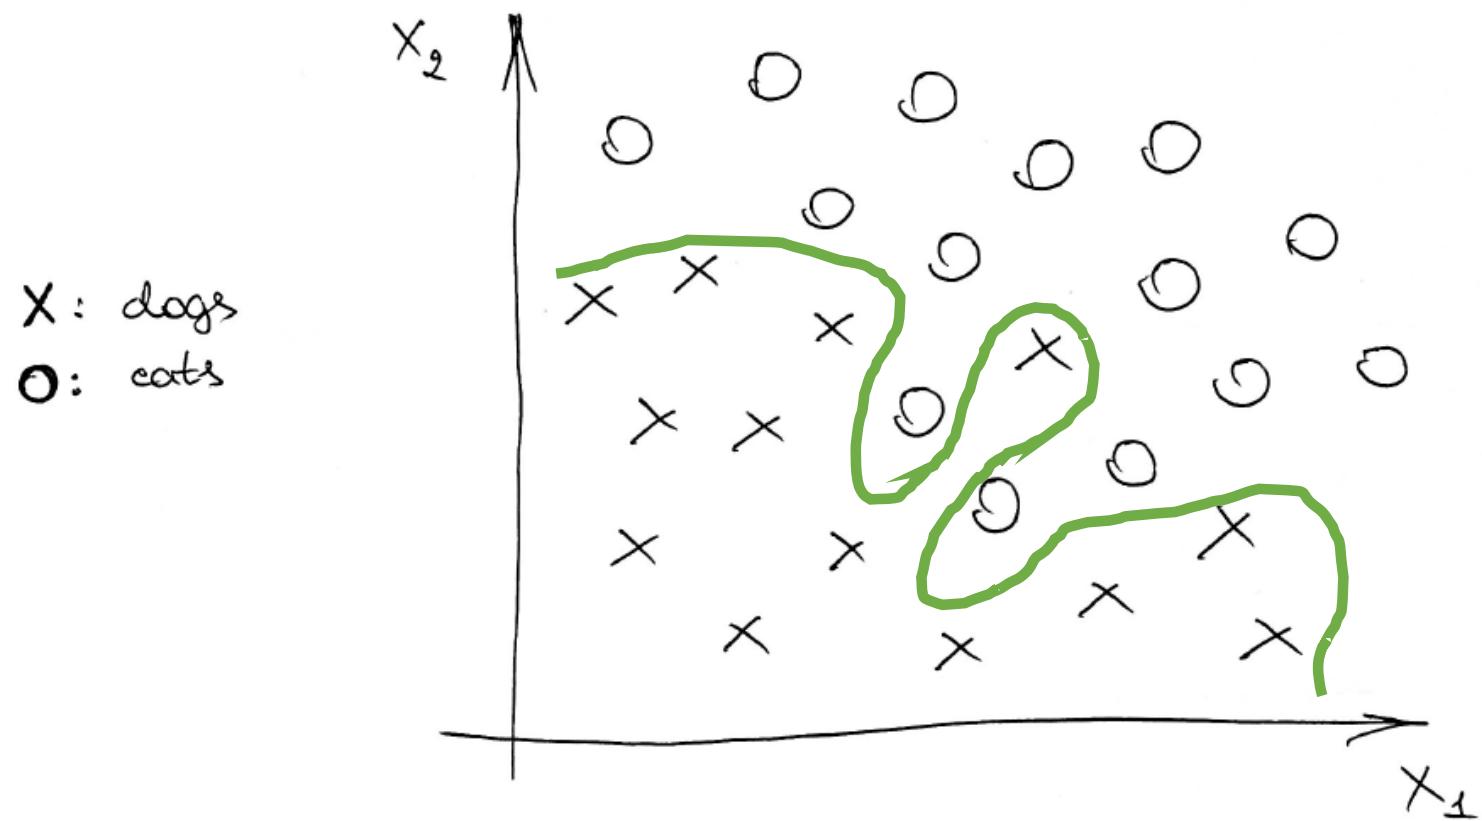
$$W = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

If $-1 + x_1^2 + x_2^2 \geq 0 \Rightarrow y = 0''$

Overfitting

$$h_w(x) = g(W_0 + W_1 x_1 + W_2 x_2 + \dots + W_{100000} x_1 + \dots)$$

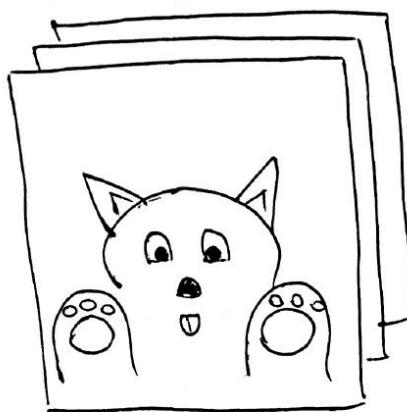
Overfitting



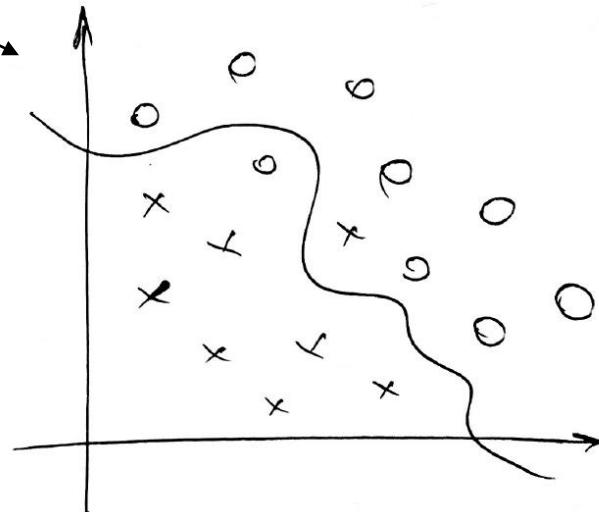
Mathematical framework

- Program is said to learn from **experience E** with respect to some class of **tasks T** and performance **measure P** if performance improves

Mitchell, 1997



$$\{ R^{250 \times 250} \}$$



%

correct
answers
 $\{ \text{cat, dog} \}$

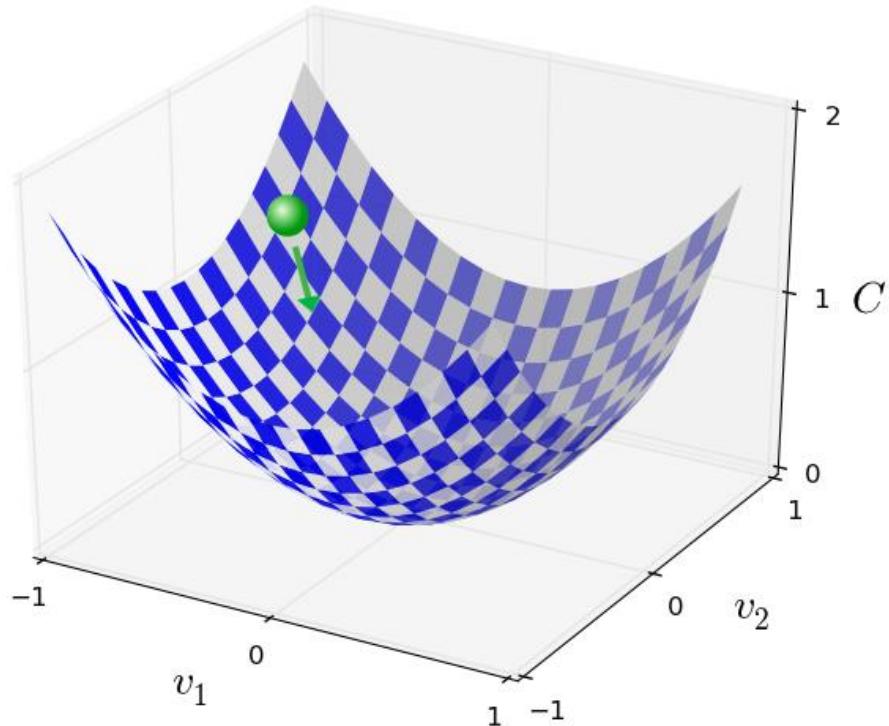
$\rightarrow \max$

$$f: R^{250 \times 250} \rightarrow \{ \text{cat, dog} \}$$

Measure

$$\frac{1}{2m} \sum_{i=1}^m \left(h_w(x^{(i)}) - y^{(i)} \right)^2$$

How it learns?



- We want to minimize some $J(x_1, x_2)$
- **Gradient (∇J)** – vector, showing the direction of the fastest **increasing** of J .
- To minimize, we have to move in the direction of **antigradient** $-\nabla C$ with some **step** λ .

Gradient descent

$$y(x_1, x_2) = x_1^2 + x_2^2$$

$$x^{(0)} = (5, 5)^T$$

$$\lambda = 0.1$$

UPDATE RULE

$$x^{(k+1)} = x^{(k)} - \lambda \cdot \nabla y(x^{(k)})$$

Gradient descent

$$\nabla J(x^{(0)}) = \begin{pmatrix} 2 \cdot x_1^{(0)} \\ 2 \cdot x_2^{(0)} \end{pmatrix} = \begin{pmatrix} 10 \\ 10 \end{pmatrix}$$

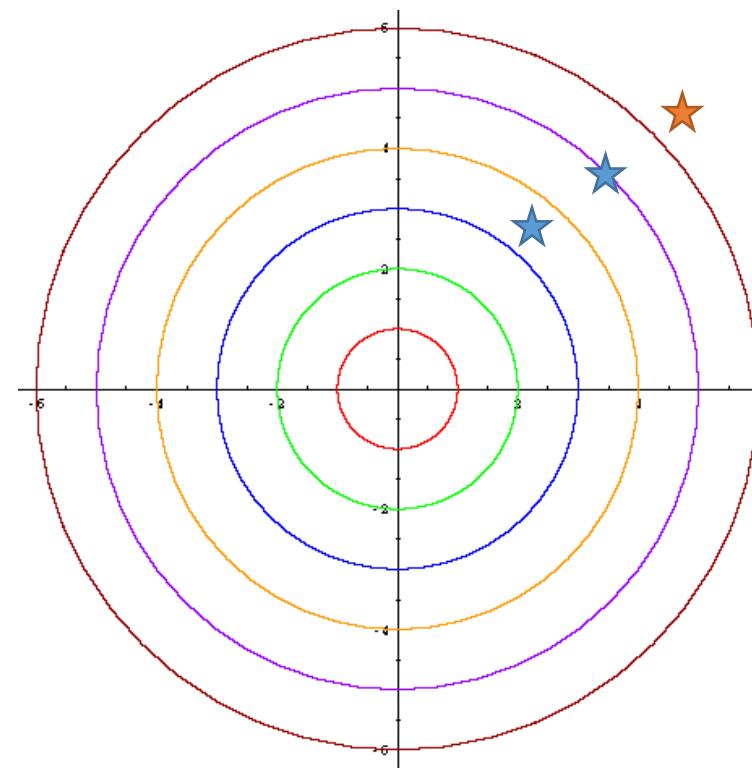
$$\begin{aligned} x^{(1)} &= \begin{pmatrix} 5 \\ 5 \end{pmatrix} - \lambda \cdot \begin{pmatrix} 10 \\ 10 \end{pmatrix} = \begin{pmatrix} 5 \\ 5 \end{pmatrix} - 0.1 \cdot \begin{pmatrix} 10 \\ 10 \end{pmatrix} = \\ &= \begin{pmatrix} 5 \\ 5 \end{pmatrix} - \begin{pmatrix} 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 4 \\ 4 \end{pmatrix} \end{aligned}$$

Gradient descent

$$\nabla J(x^{(1)}) = \begin{pmatrix} 2 \cdot x_1^{(1)} \\ 2 \cdot x_2^{(1)} \end{pmatrix} = \begin{pmatrix} 8 \\ 8 \end{pmatrix}$$

$$\begin{aligned}x^{(2)} &= \begin{pmatrix} 4 \\ 4 \end{pmatrix} - \lambda \cdot \begin{pmatrix} 8 \\ 8 \end{pmatrix} = \begin{pmatrix} 4 \\ 4 \end{pmatrix} - 0.1 \cdot \begin{pmatrix} 8 \\ 8 \end{pmatrix} = \\&= \begin{pmatrix} 4 \\ 4 \end{pmatrix} - \begin{pmatrix} 0.8 \\ 0.8 \end{pmatrix} = \begin{pmatrix} 3.2 \\ 3.2 \end{pmatrix}\end{aligned}$$

Gradient descent



Gradient descent

$$h_w(x_1, x_2) = W_0 + W_1 x_1 + W_2 x_2$$

$$\mathcal{J}(W_0, W_1, W_2) = \frac{1}{2m} \sum_{i=1}^m (h_w(x_1^{(i)}, x_2^{(i)}) - y^{(i)})^2$$

$$\begin{aligned}\frac{\partial}{\partial W_0} \mathcal{J}(W_0, W_1, W_2) &= \frac{\partial}{\partial W_0} \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2 = \\&= \frac{\partial}{\partial W_0} \frac{1}{2m} \sum_{i=1}^m ((W_0 + W_1 x_1^{(i)} + W_2 x_2^{(i)})^2 - 2y^{(i)}(W_0 + W_1 x_1^{(i)} + W_2 x_2^{(i)}) + y^2) = \\&= \frac{1}{2m} \sum_{i=1}^m (2(W_0 + W_1 x_1^{(i)} + W_2 x_2^{(i)}) - 2y^{(i)}) = \\&= \frac{1}{m} \sum_{i=1}^m (W_0 + W_1 x_1^{(i)} + W_2 x_2^{(i)} - y^{(i)}) = \\&= \frac{1}{m} \sum_{i=1}^m (h_w(x_1^{(i)}, x_2^{(i)}) - y^{(i)})\end{aligned}$$

Gradient descent

$$\frac{\partial}{\partial w_0} J(\bar{w}) = \frac{1}{m} \sum_{i=1}^m (h_w(\bar{x}^{(i)}) - y^{(i)})$$

$$\frac{\partial}{\partial w_1} J(\bar{w}) = \frac{1}{m} \sum_{i=1}^m (h_w(\bar{x}^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$\frac{\partial}{\partial w_2} J(\bar{w}) = \frac{1}{m} \sum_{i=1}^m (h_w(\bar{x}^{(i)}) - y^{(i)}) x_2^{(i)}$$

How it learns?

Program is said to learn from **experience E** (data set with pictures) with respect to some class of **tasks T** (logistic regression) and performance **measure P** (mean squared error) if performance improves (with gradient descent iterations)

Mitchell, 1997

deep learning in nutshell

Image representation

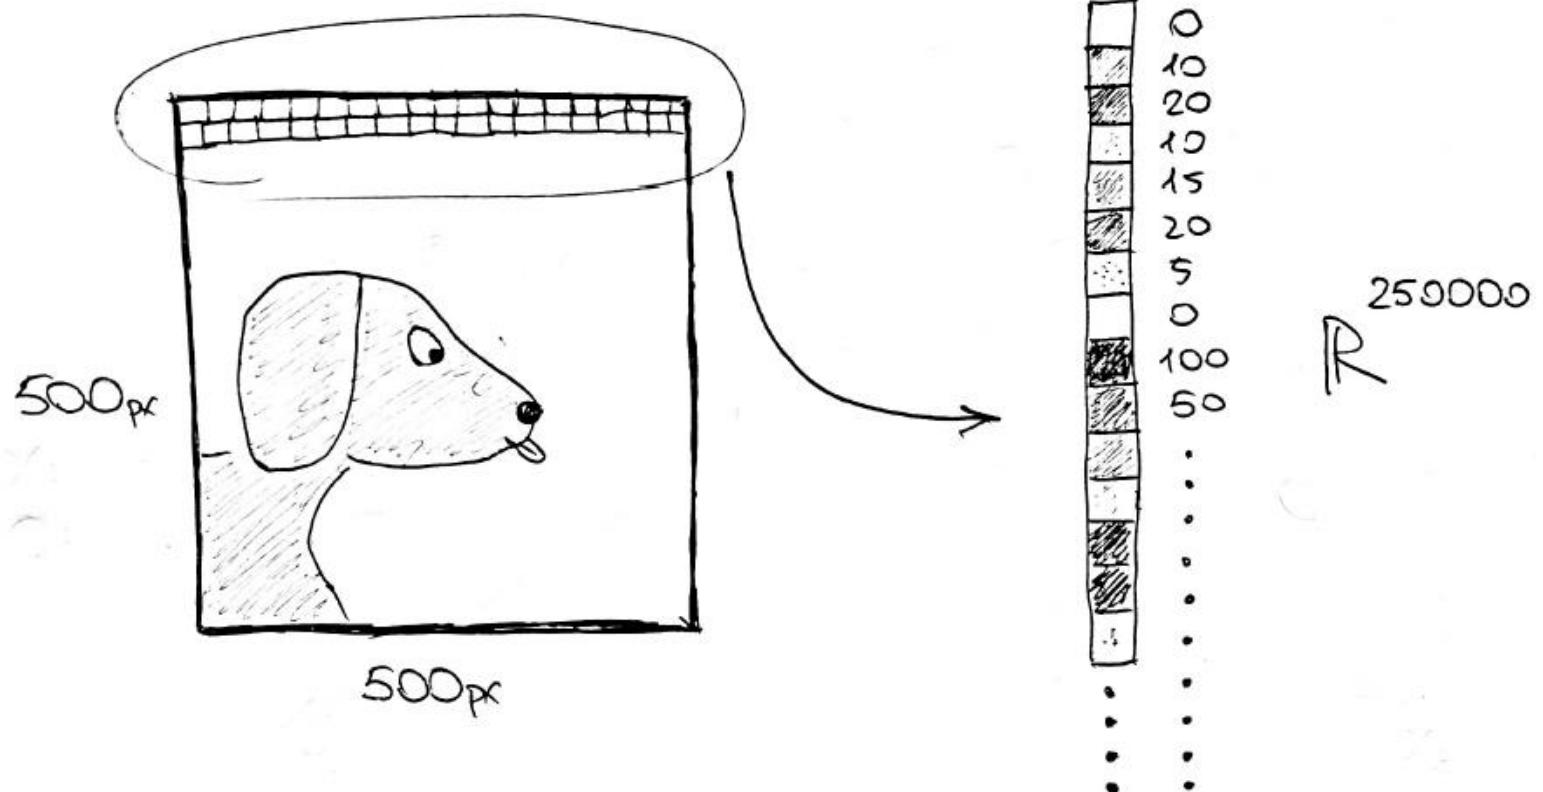


Image representation

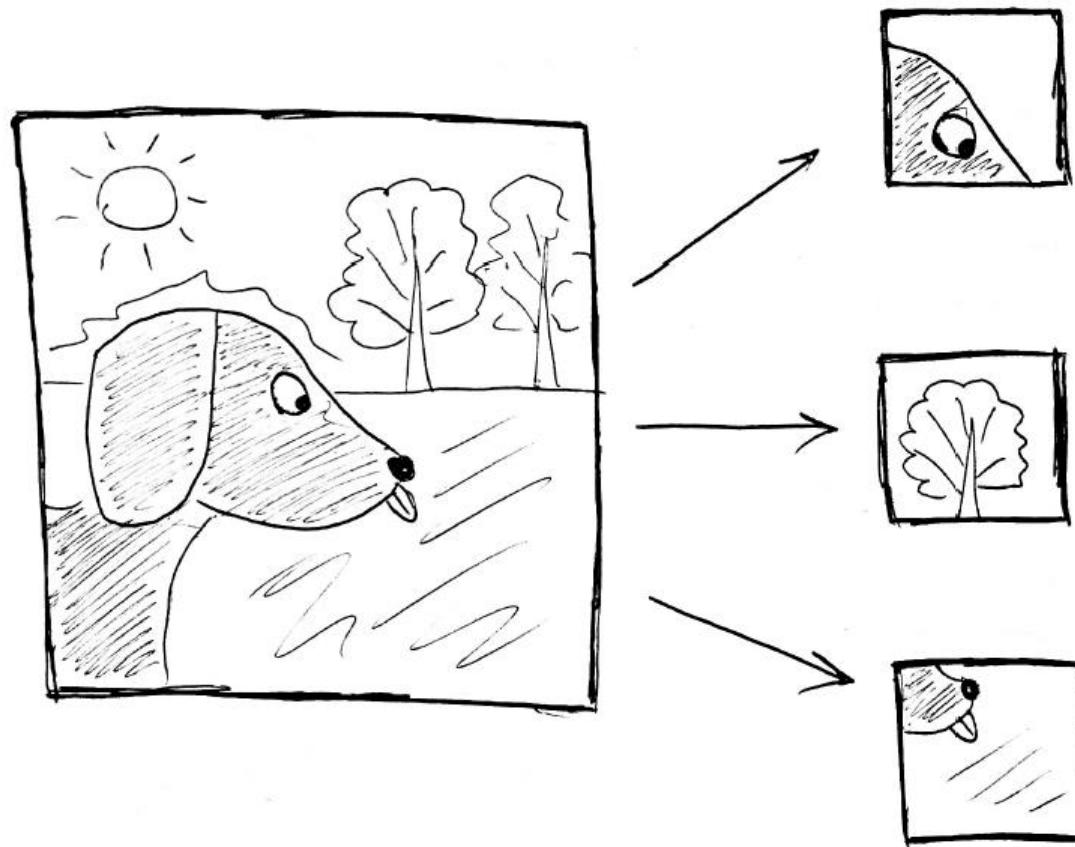


Image representation

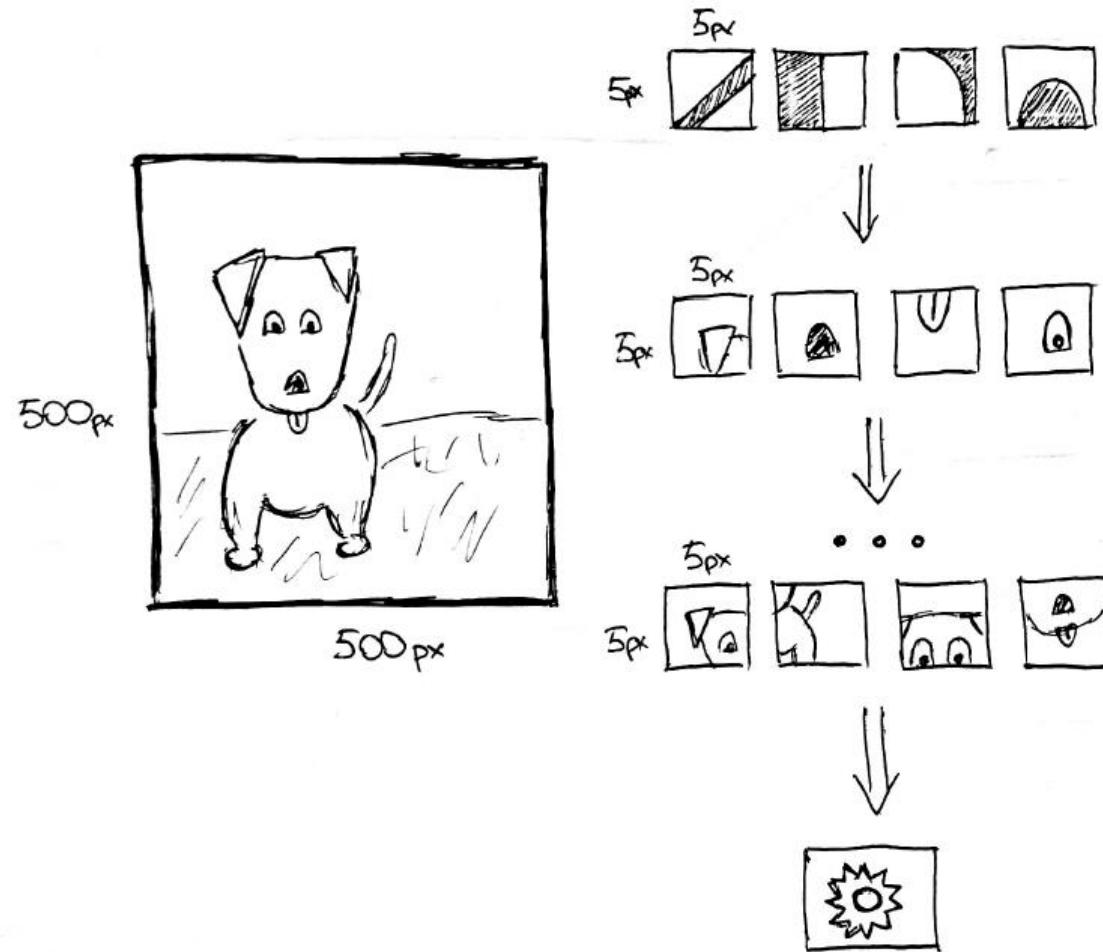


Image representation

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Image representation

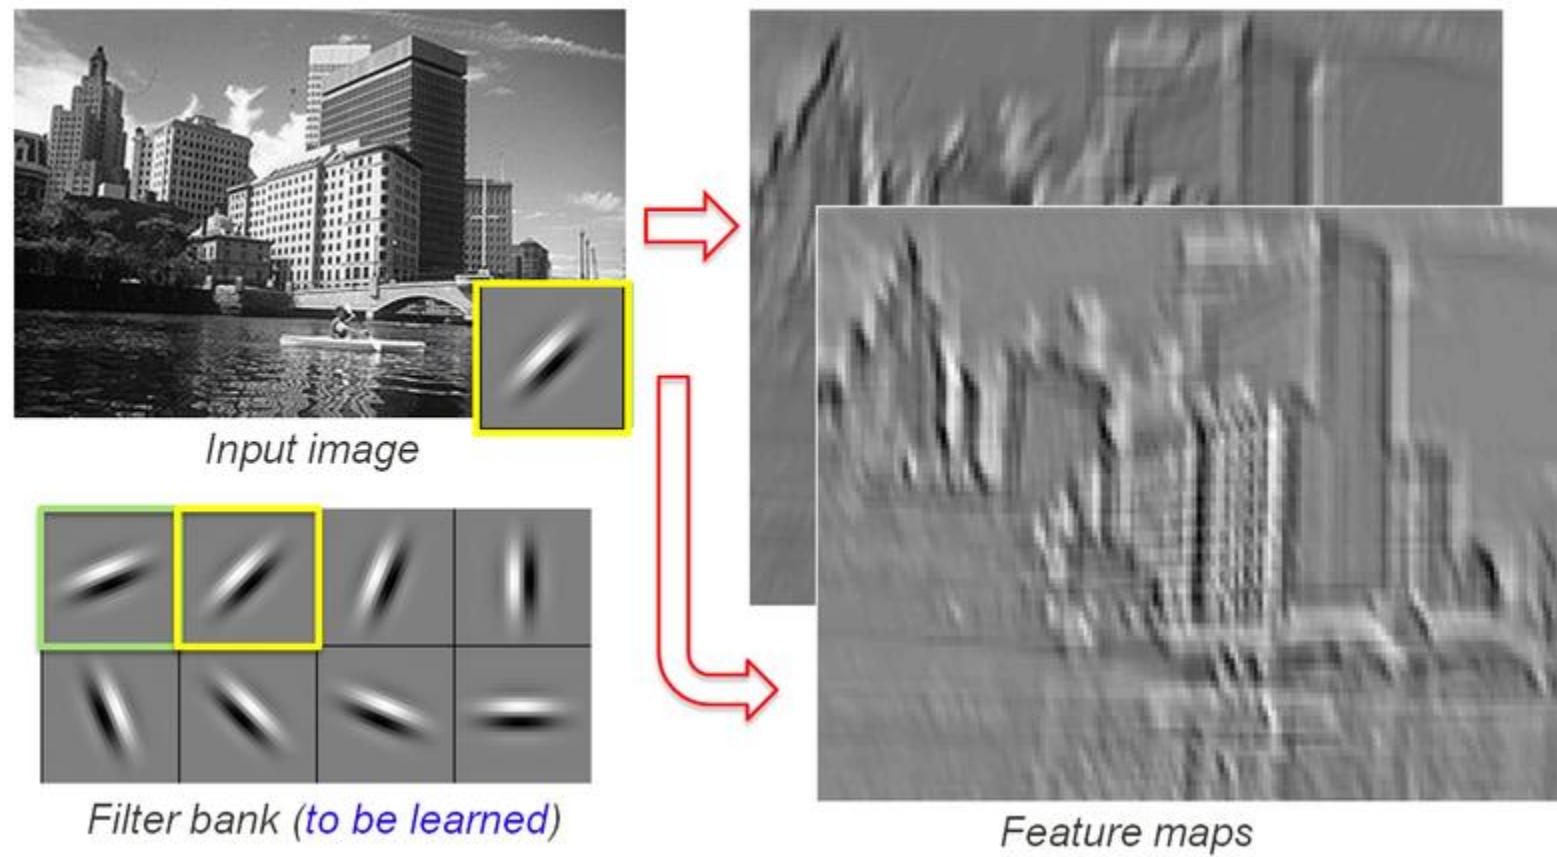
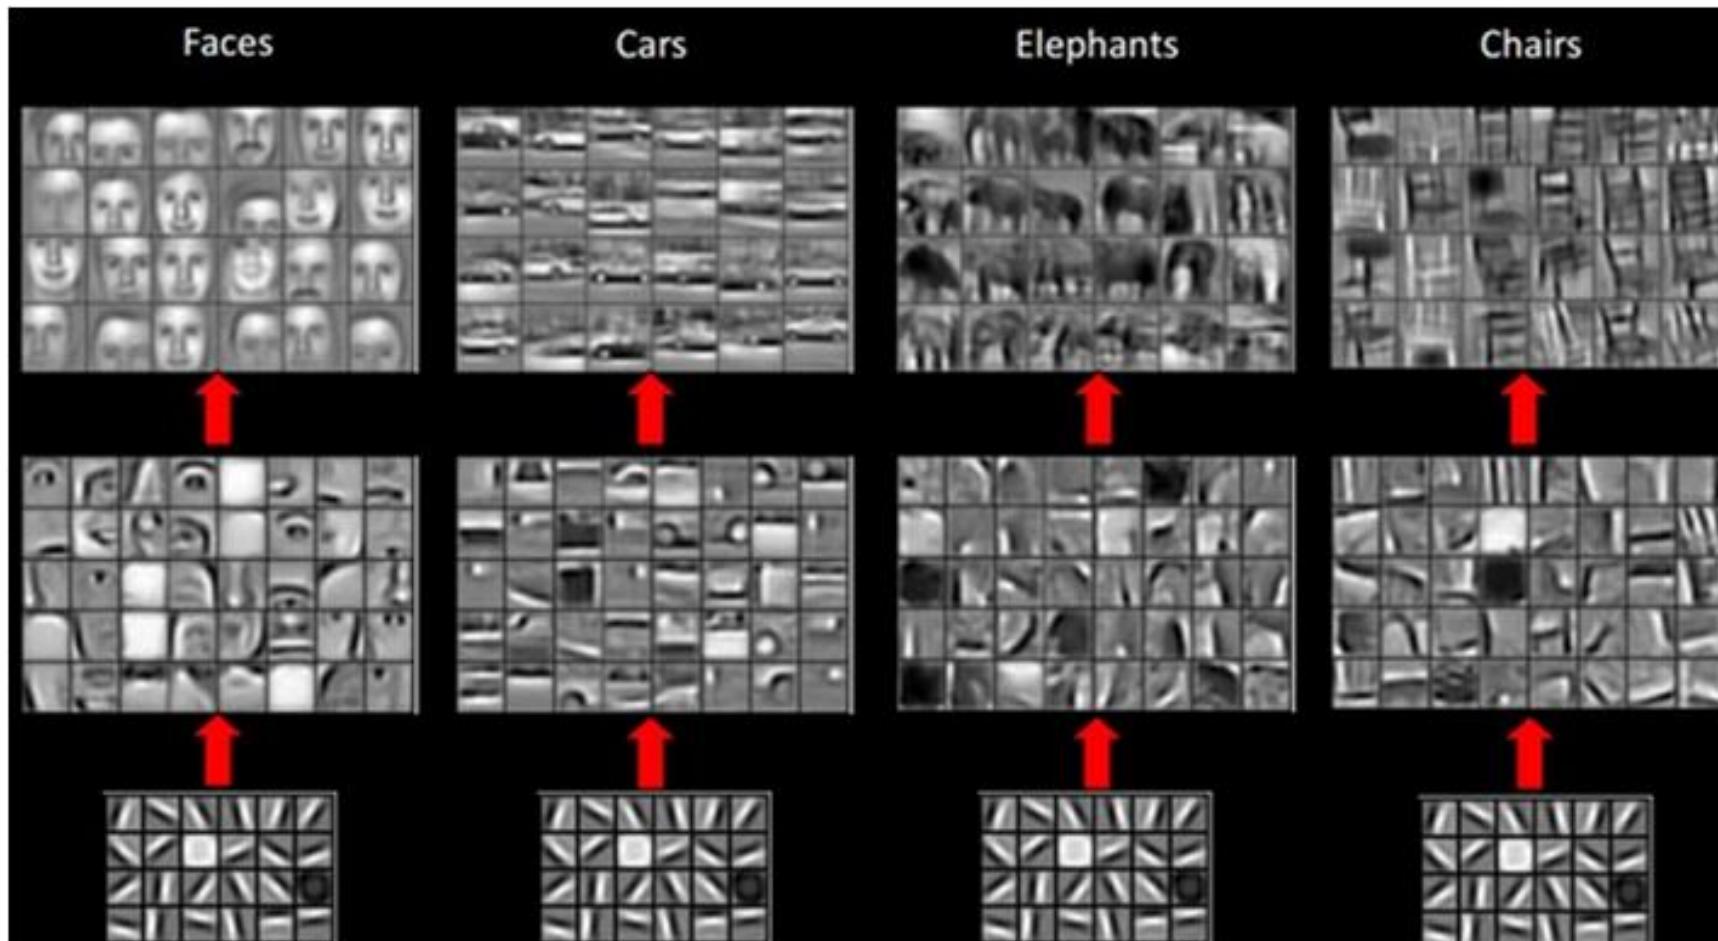
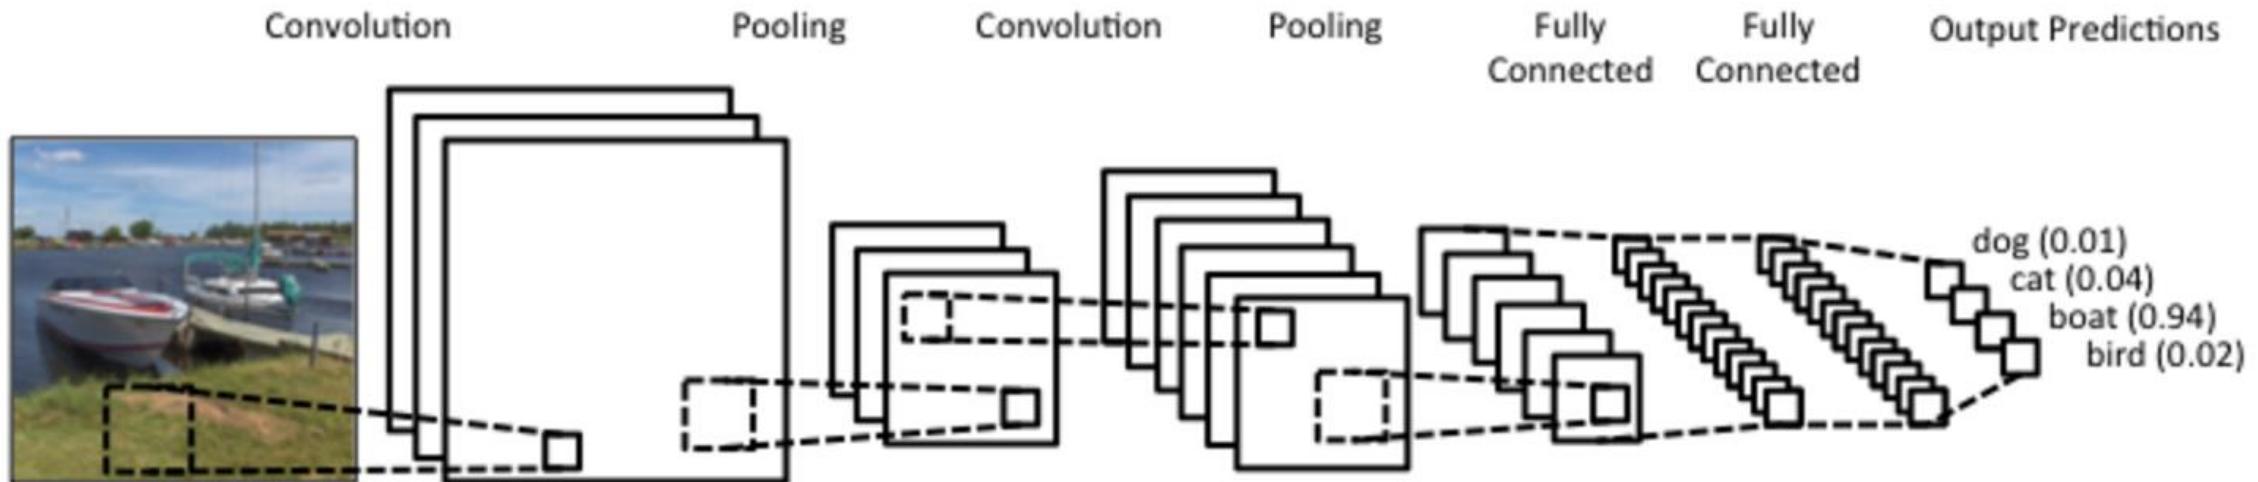


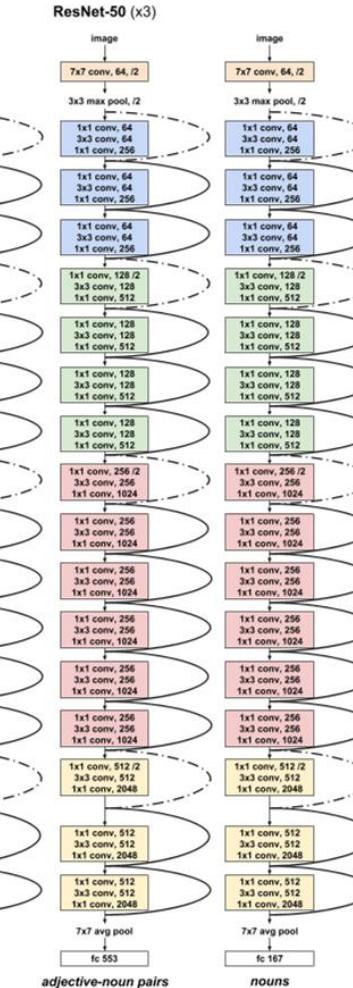
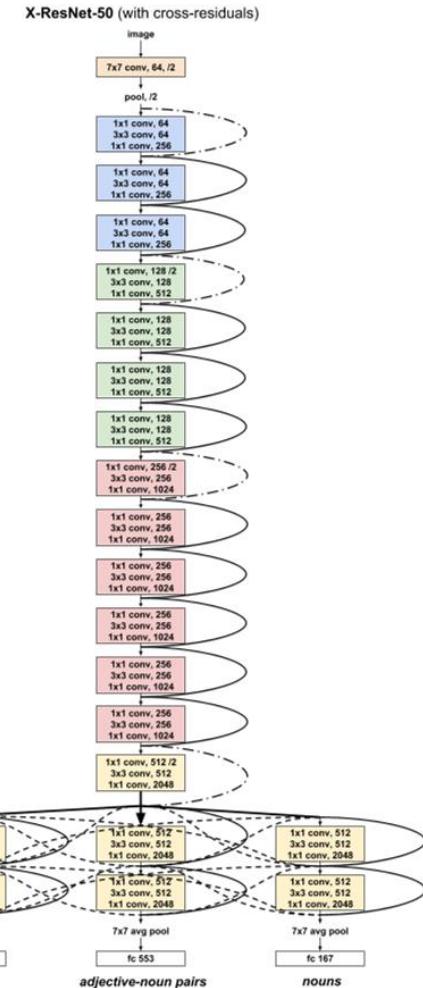
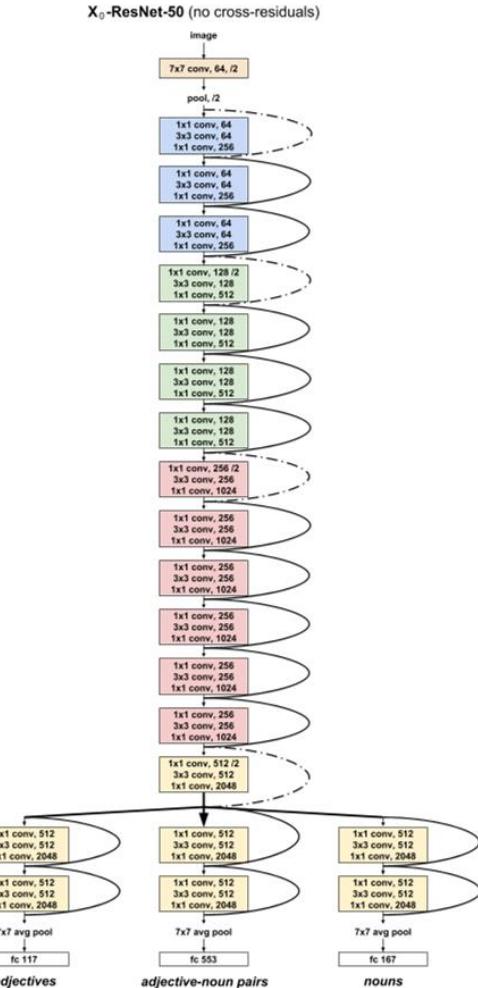
Image representation



Convolutional neural network



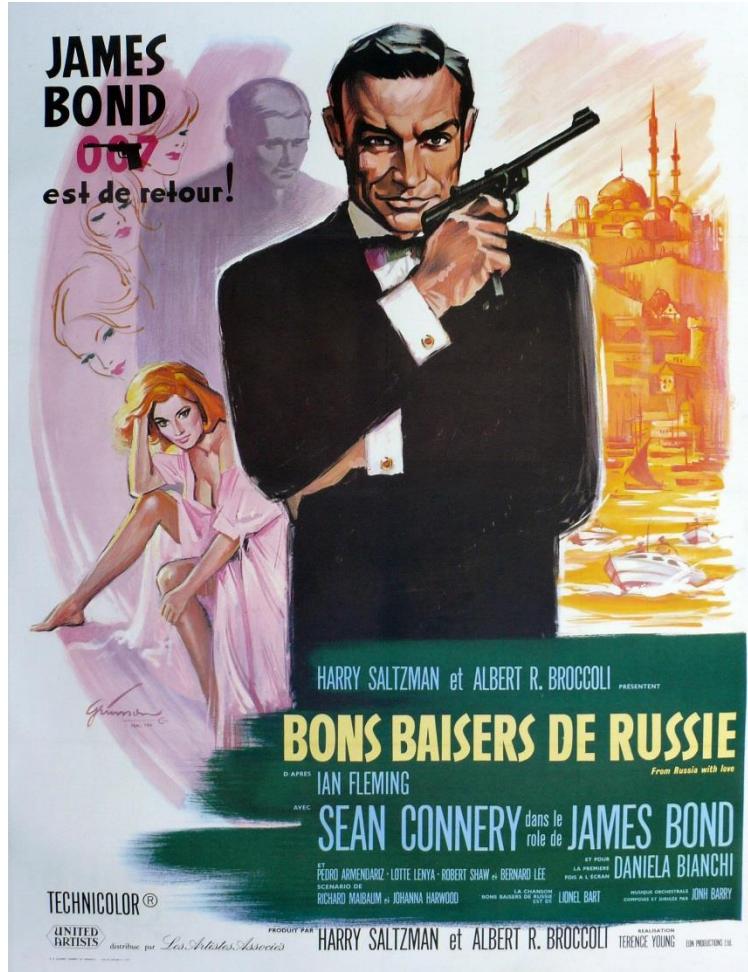
Convolutional neural network



résumé

We want to extract the most **meaningful** information from the images to pass them to some algorithm that can correctly split these representations into different categories

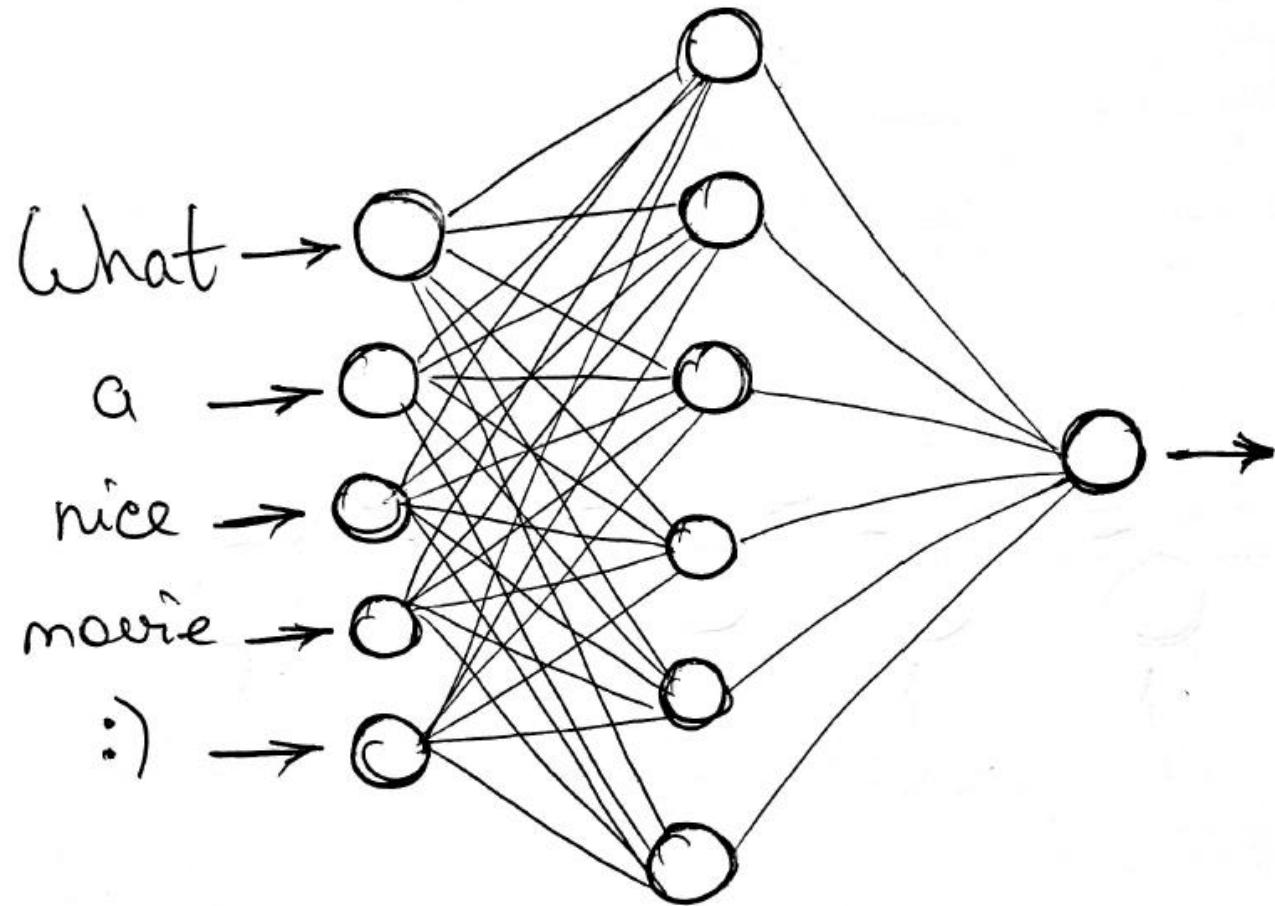
Positive vs Negative comment



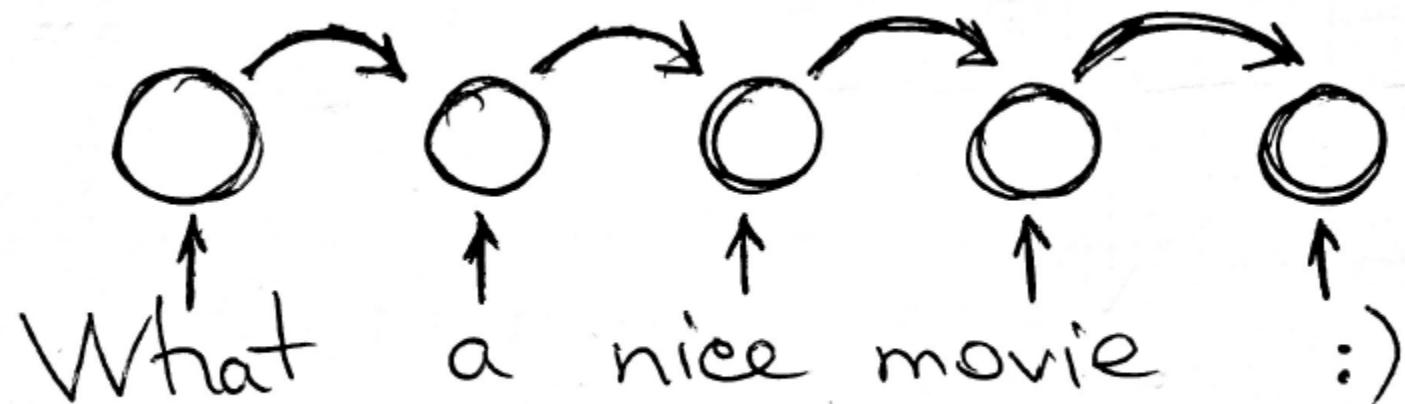
- “Amazing movie, really loved it, j’adore!!! :)"
- “Scheiss schrecklich! I wish Jason Statham played Bond, Connery is so bad :(“

Positive or Negative?

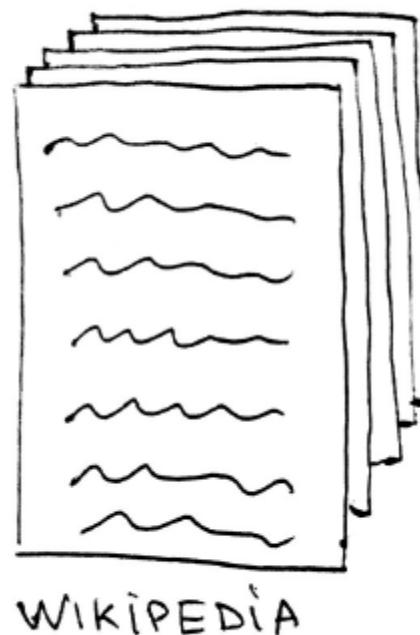
Text to neural network



Text to recurrent neural network



How to represent words?



$$\text{man} = [x_1 \dots x_{100}]$$

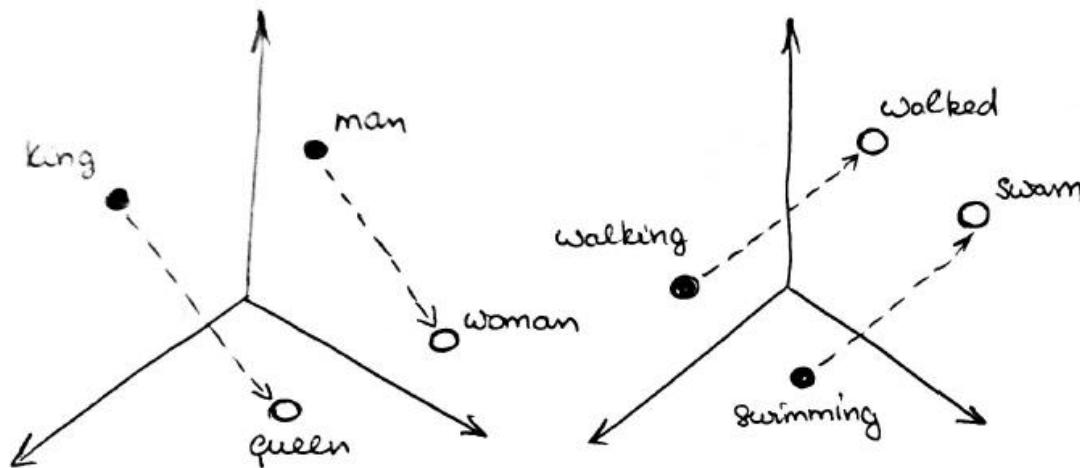
$$\text{woman} = [x_1 \dots x_{100}]$$

⋮

$$\text{king} = [x_1 \dots x_{100}]$$

$$\text{zoo} = [x_1 \dots x_{100}]$$

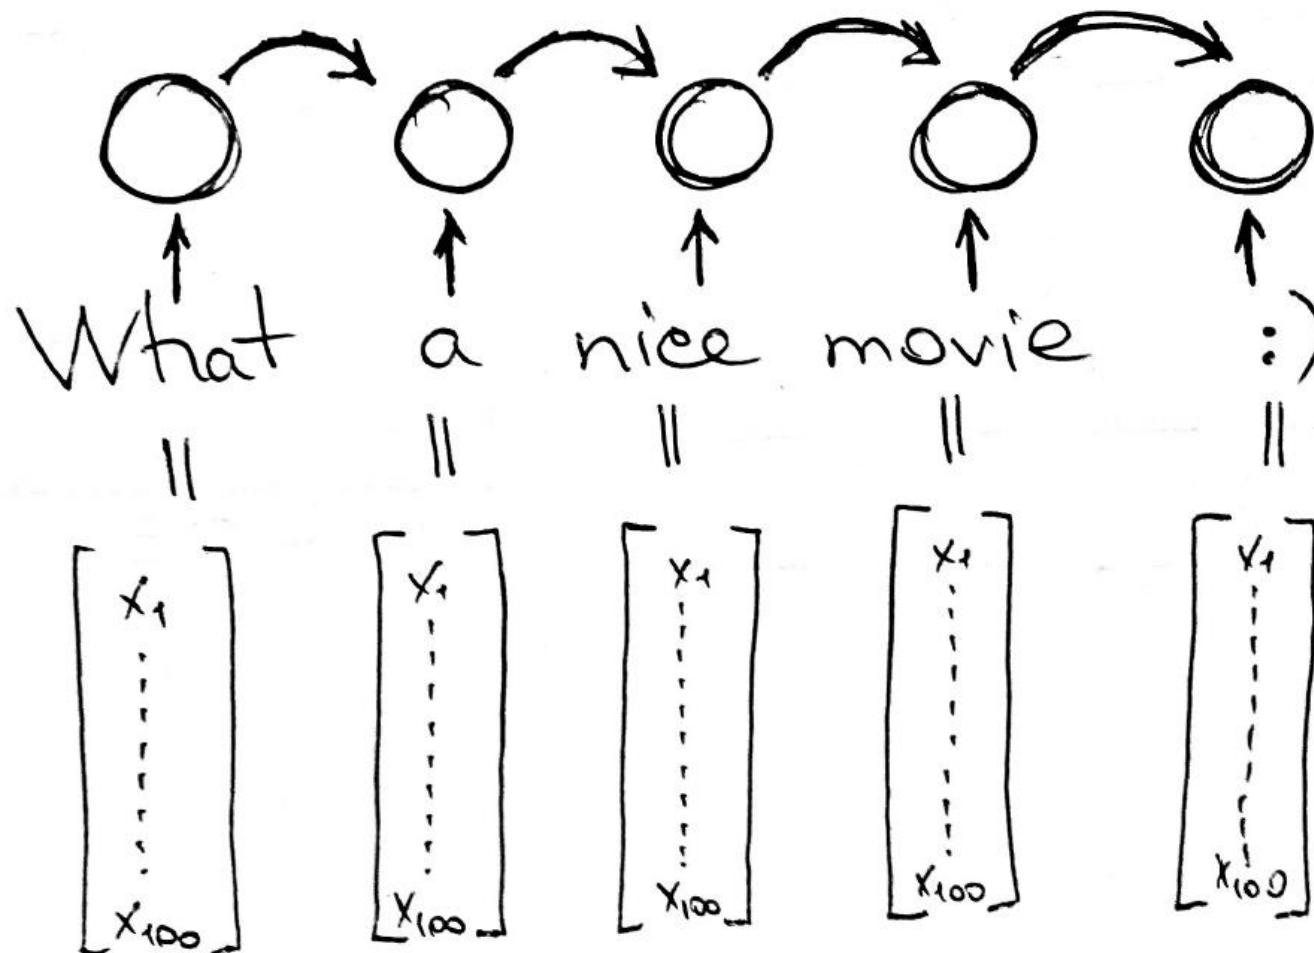
How to represent words?



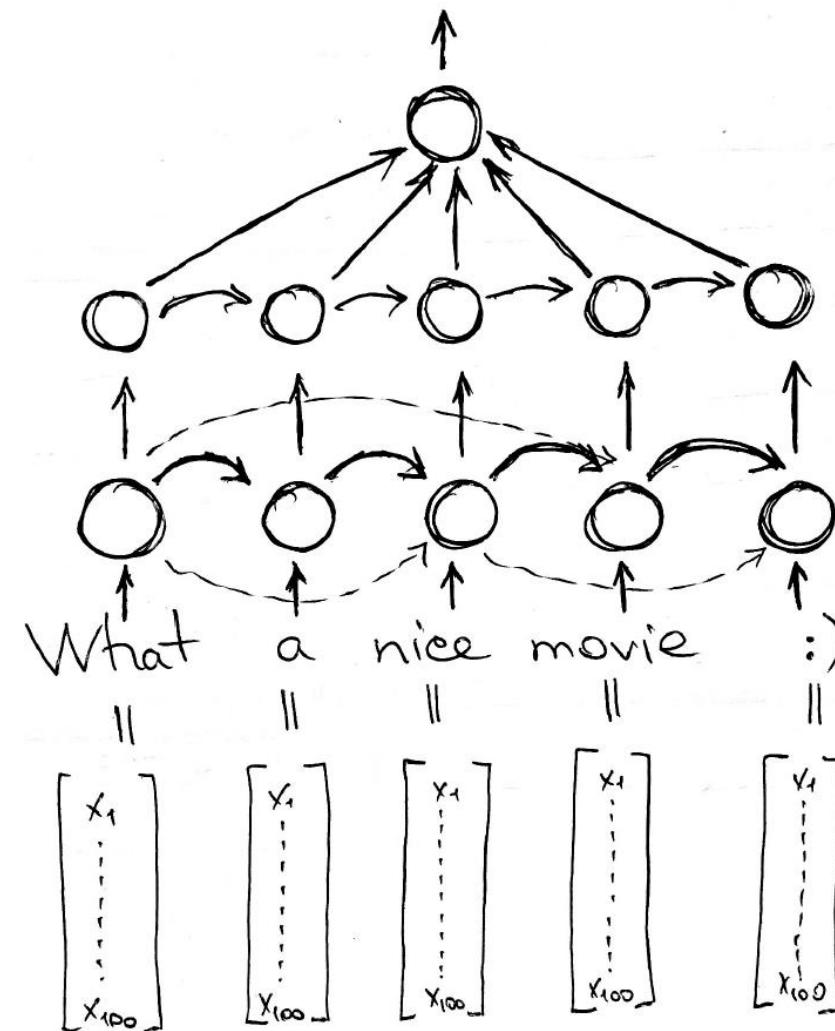
Spain ————— Madrid
Italy ————— Rome
Germany ————— Berlin
Ukraine ————— Kiev
Japan ————— Tokyo
Russia ————— Moscow
Turkey ————— Ankara

Iraq - Violence = Jordan
Human - Animal = Elephants
President - Power = Prime Minister
Library - Books = Hall

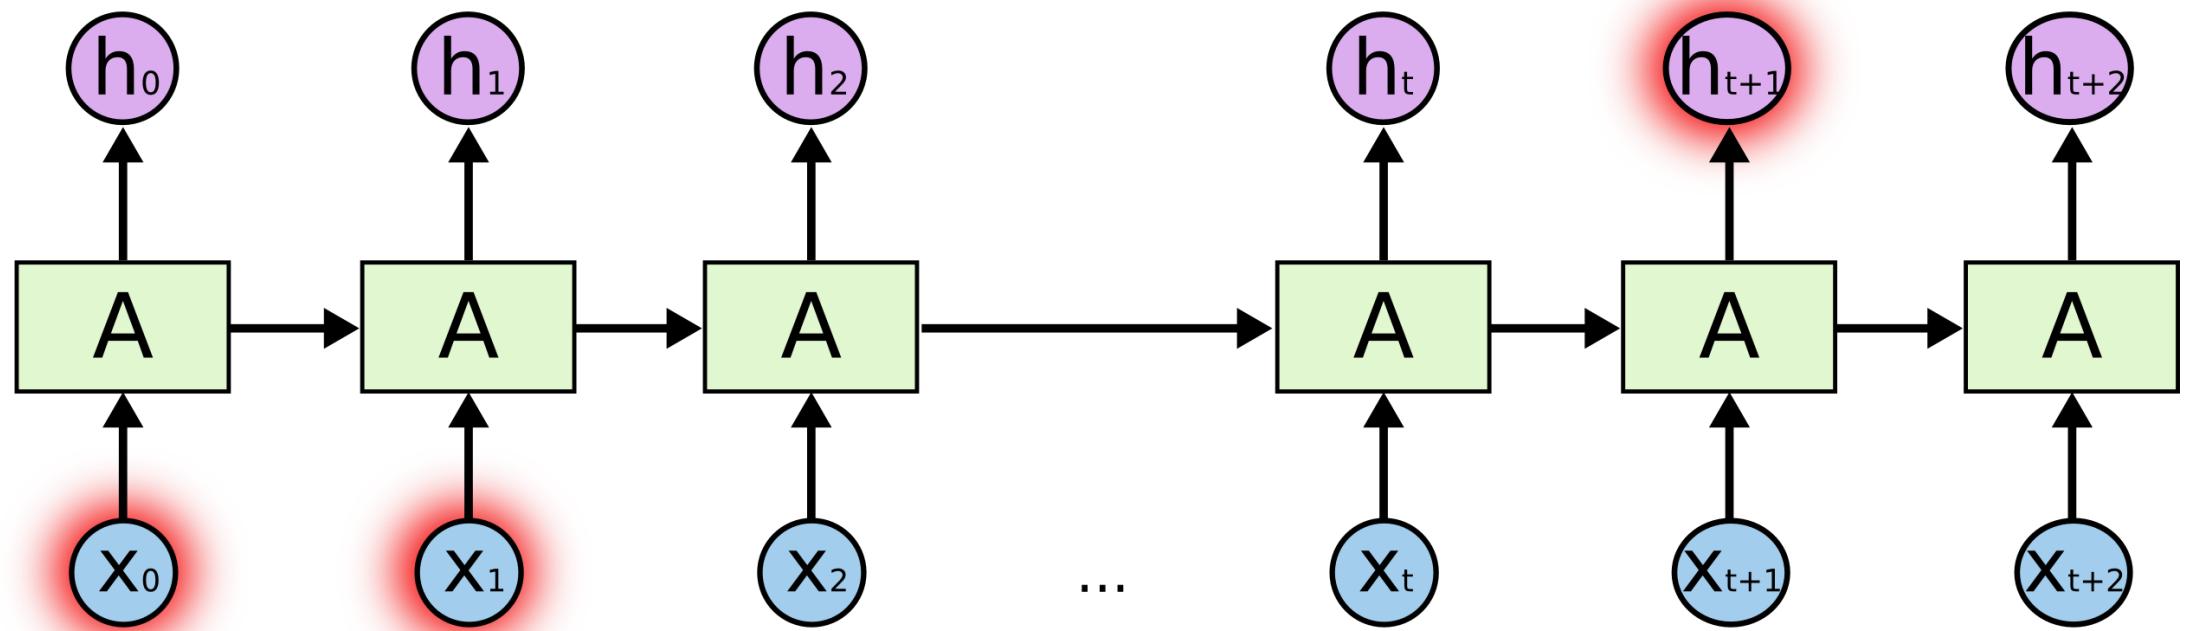
Text to recurrent neural network



Text to recurrent neural network



Recurrent neural network

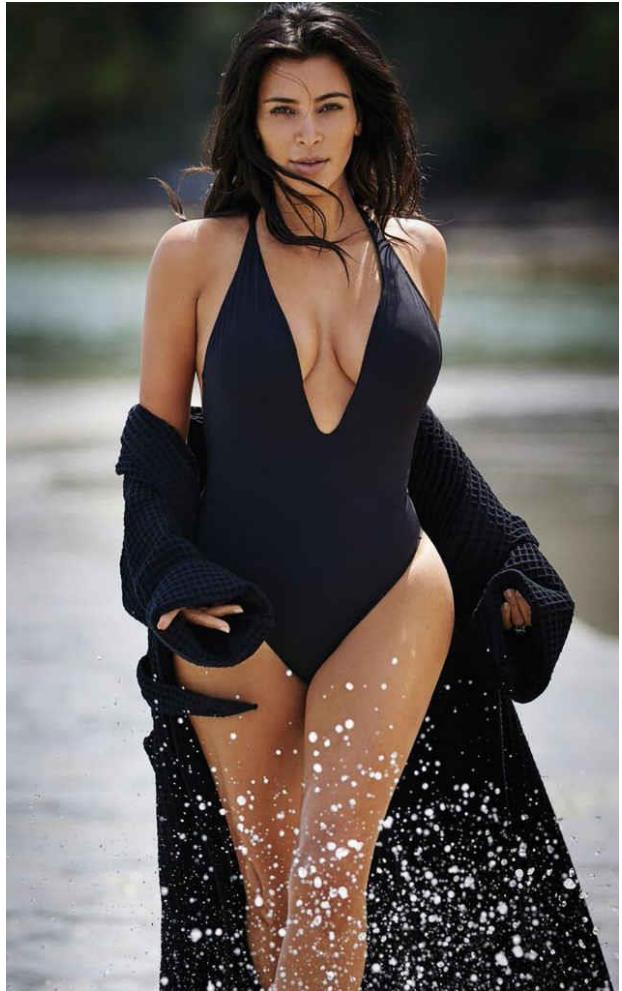


résumé

We use **recurrent** neural networks to work with **sequential** data; we also can represent all words with **vectors**, that show their sense and similarity

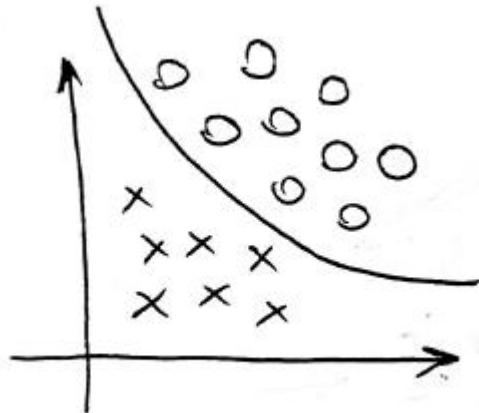
deep building blocks

Data

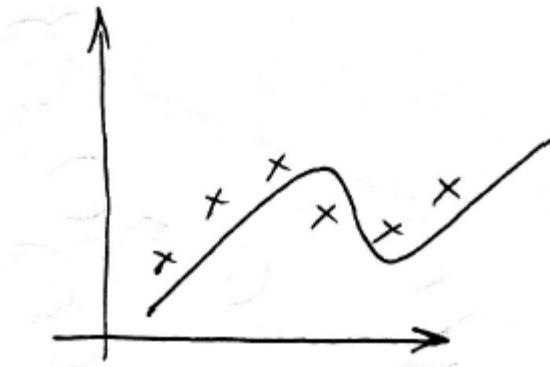


Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt. Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt.

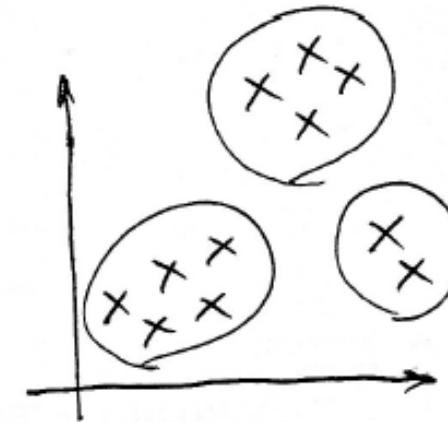
Task



Classification



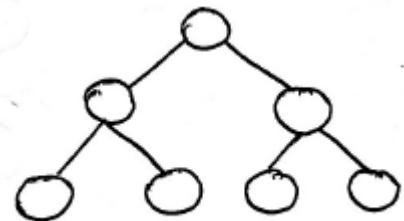
Regression (real value R)



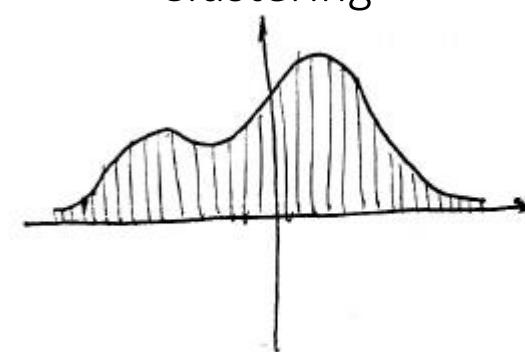
Clustering



Sequence output



Structure output (graph)

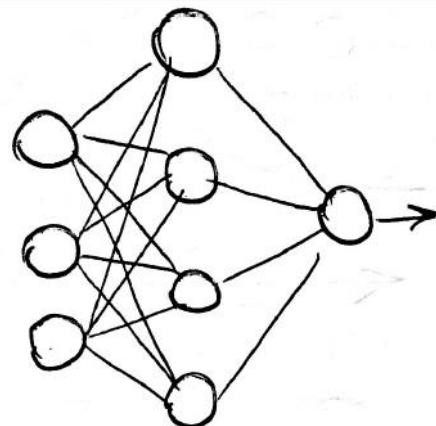


Synthesis and sampling

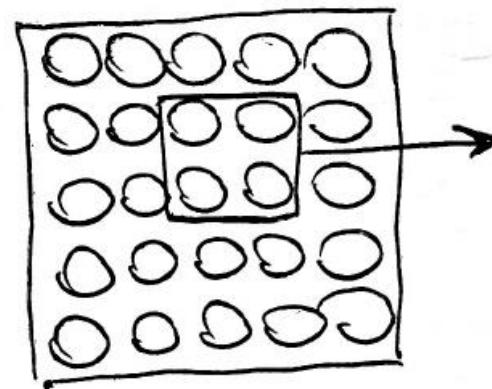
Measure (loss functions)

- Mean squared error
- Mean absolute error
- Hinge loss
- Cross-entropy loss
- Kullback-Leibler divergence
- Cosine proximity
- ...

Feature extractor



Fully-connected layer



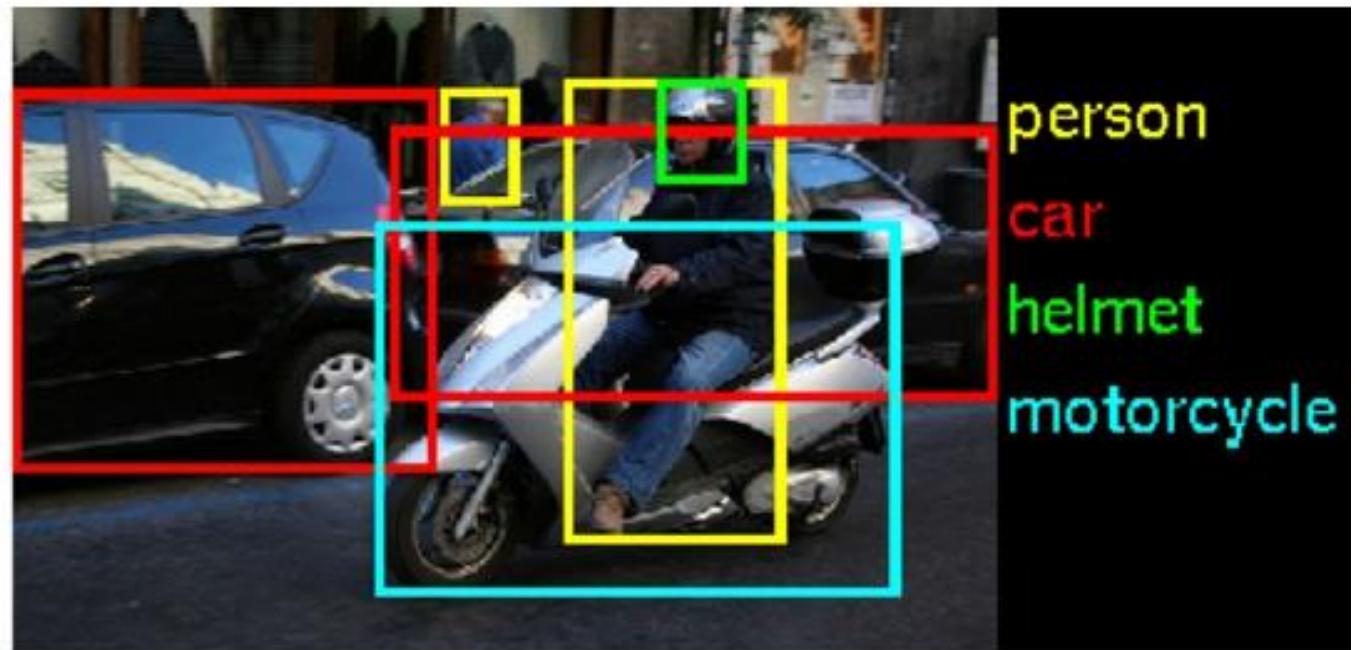
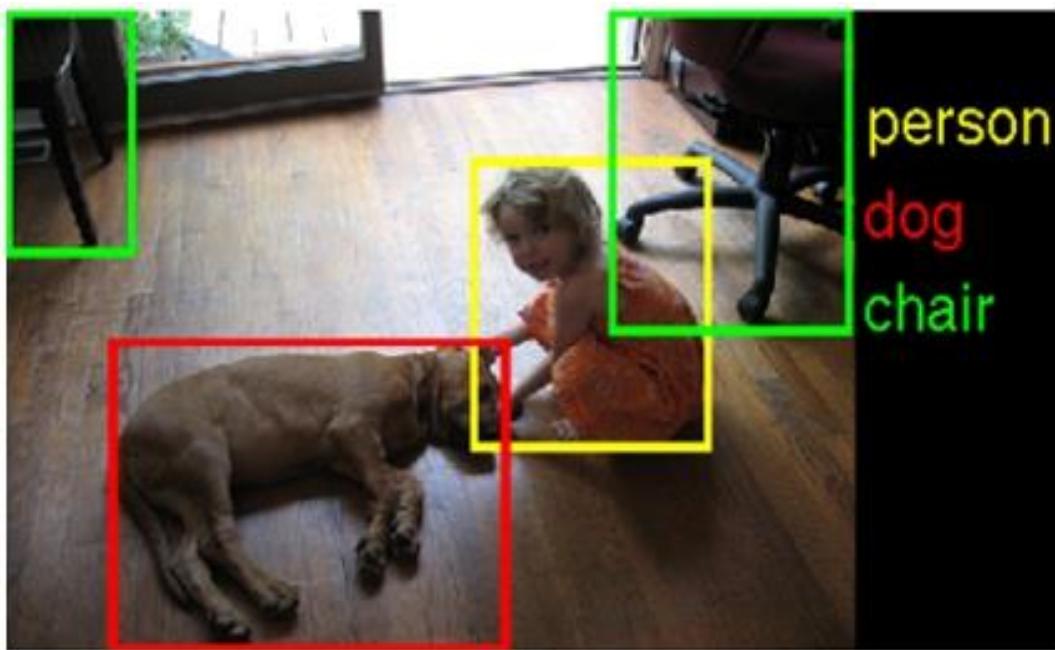
Convolutional layer



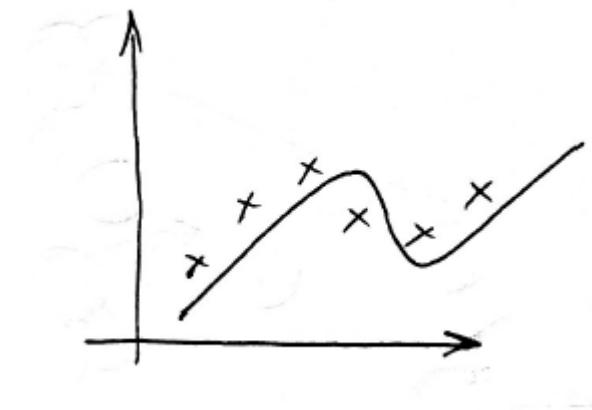
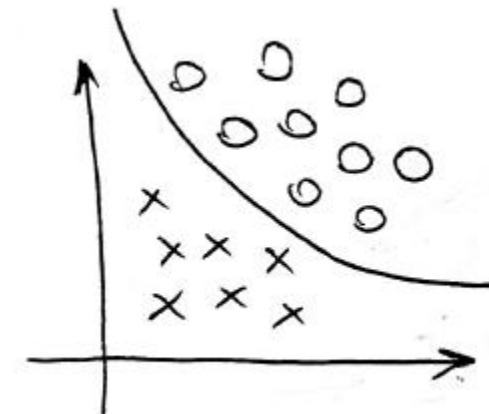
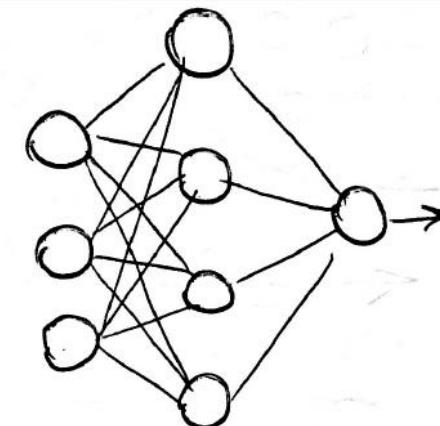
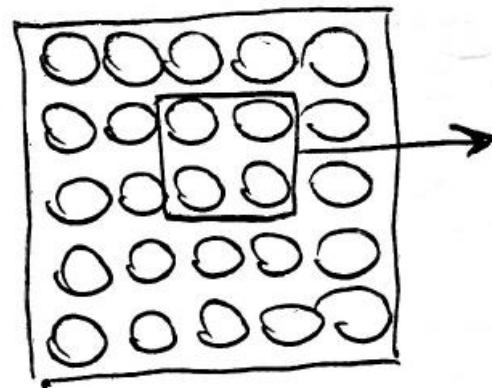
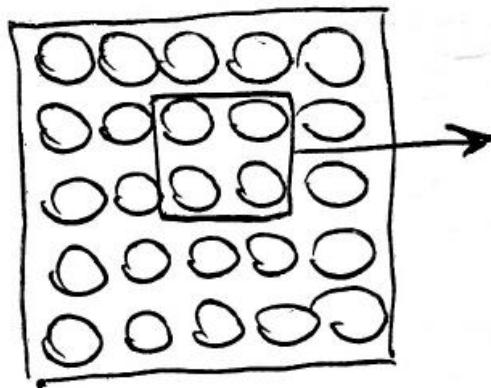
Recurrent layer

examples

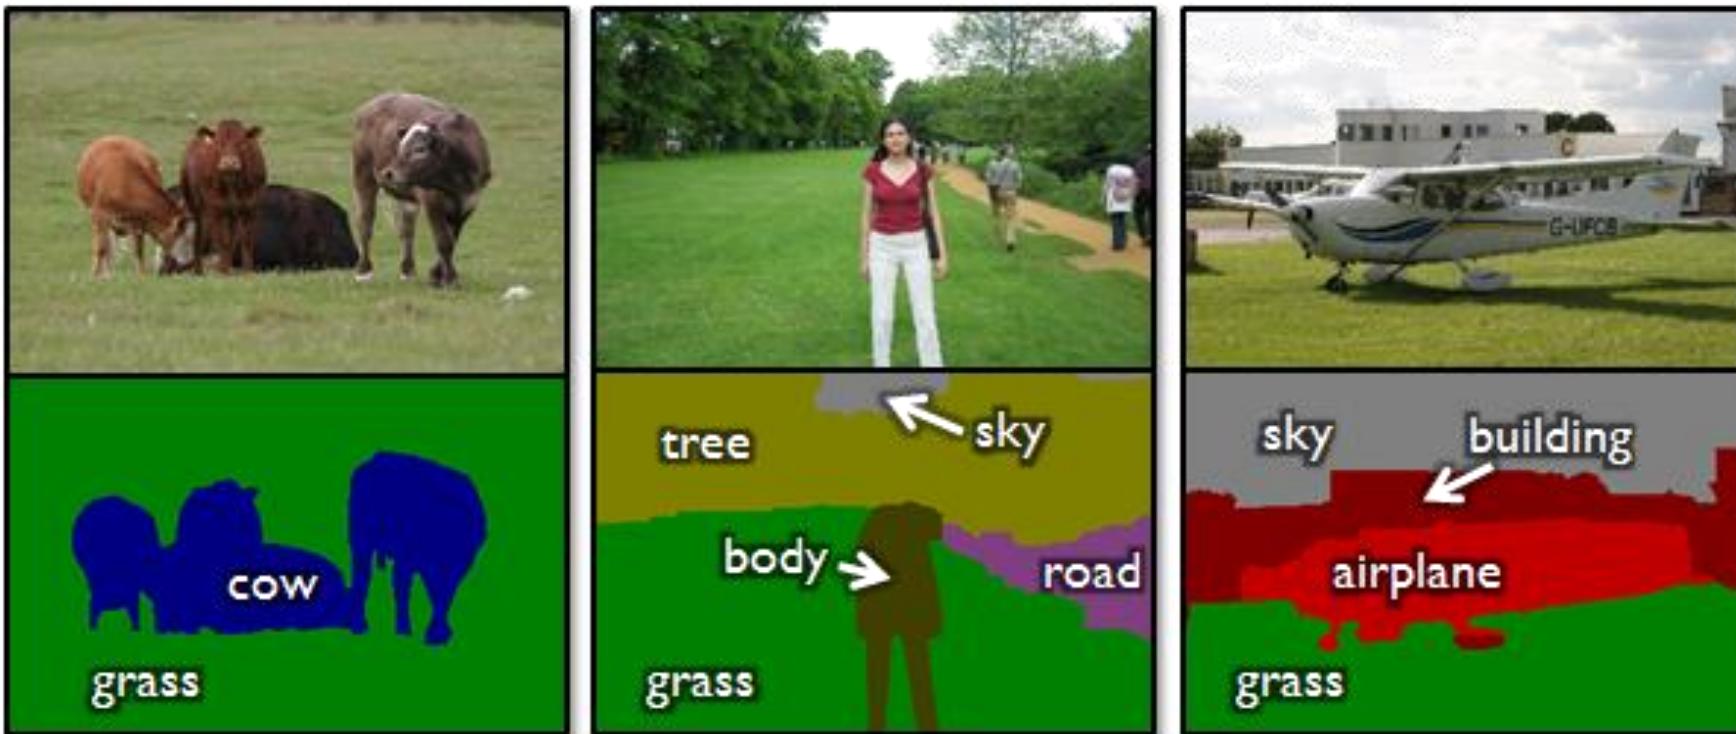
Example: Object recognition



Example: Object recognition

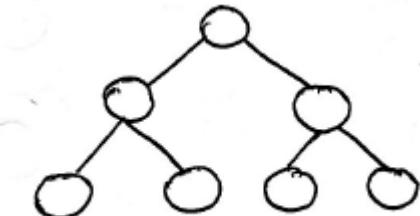
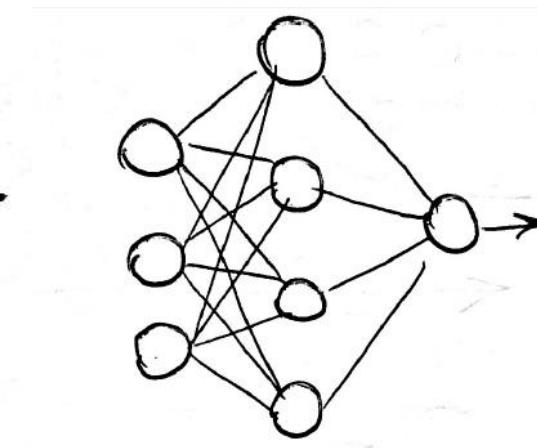
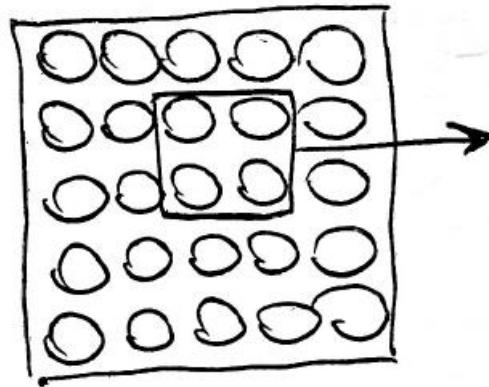
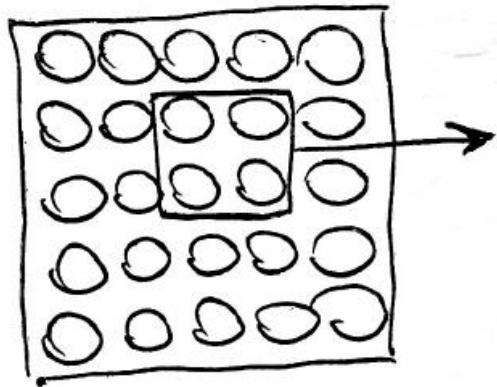


Example: Object segmentation



object classes	building	grass	tree	cow	sheep	sky	airplane	water	face	car
bicycle	flower	sign	bird	book	chair	road	cat	dog	body	boat

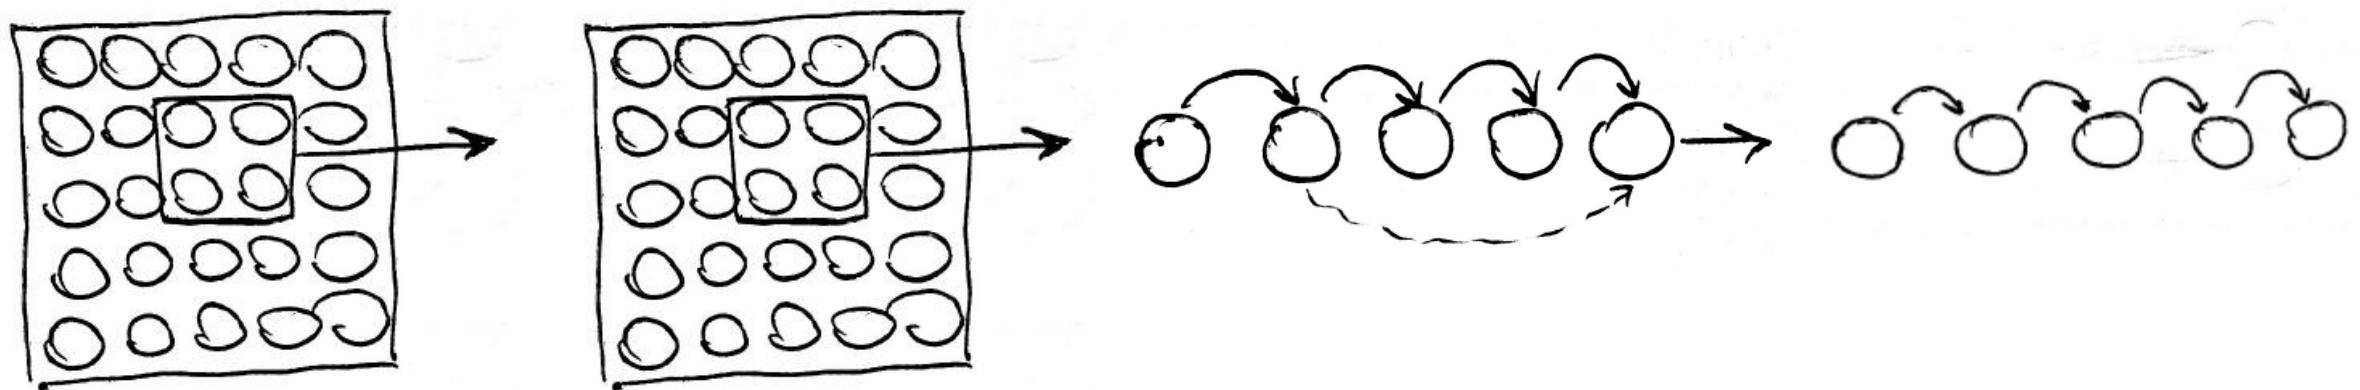
Example: Object segmentation



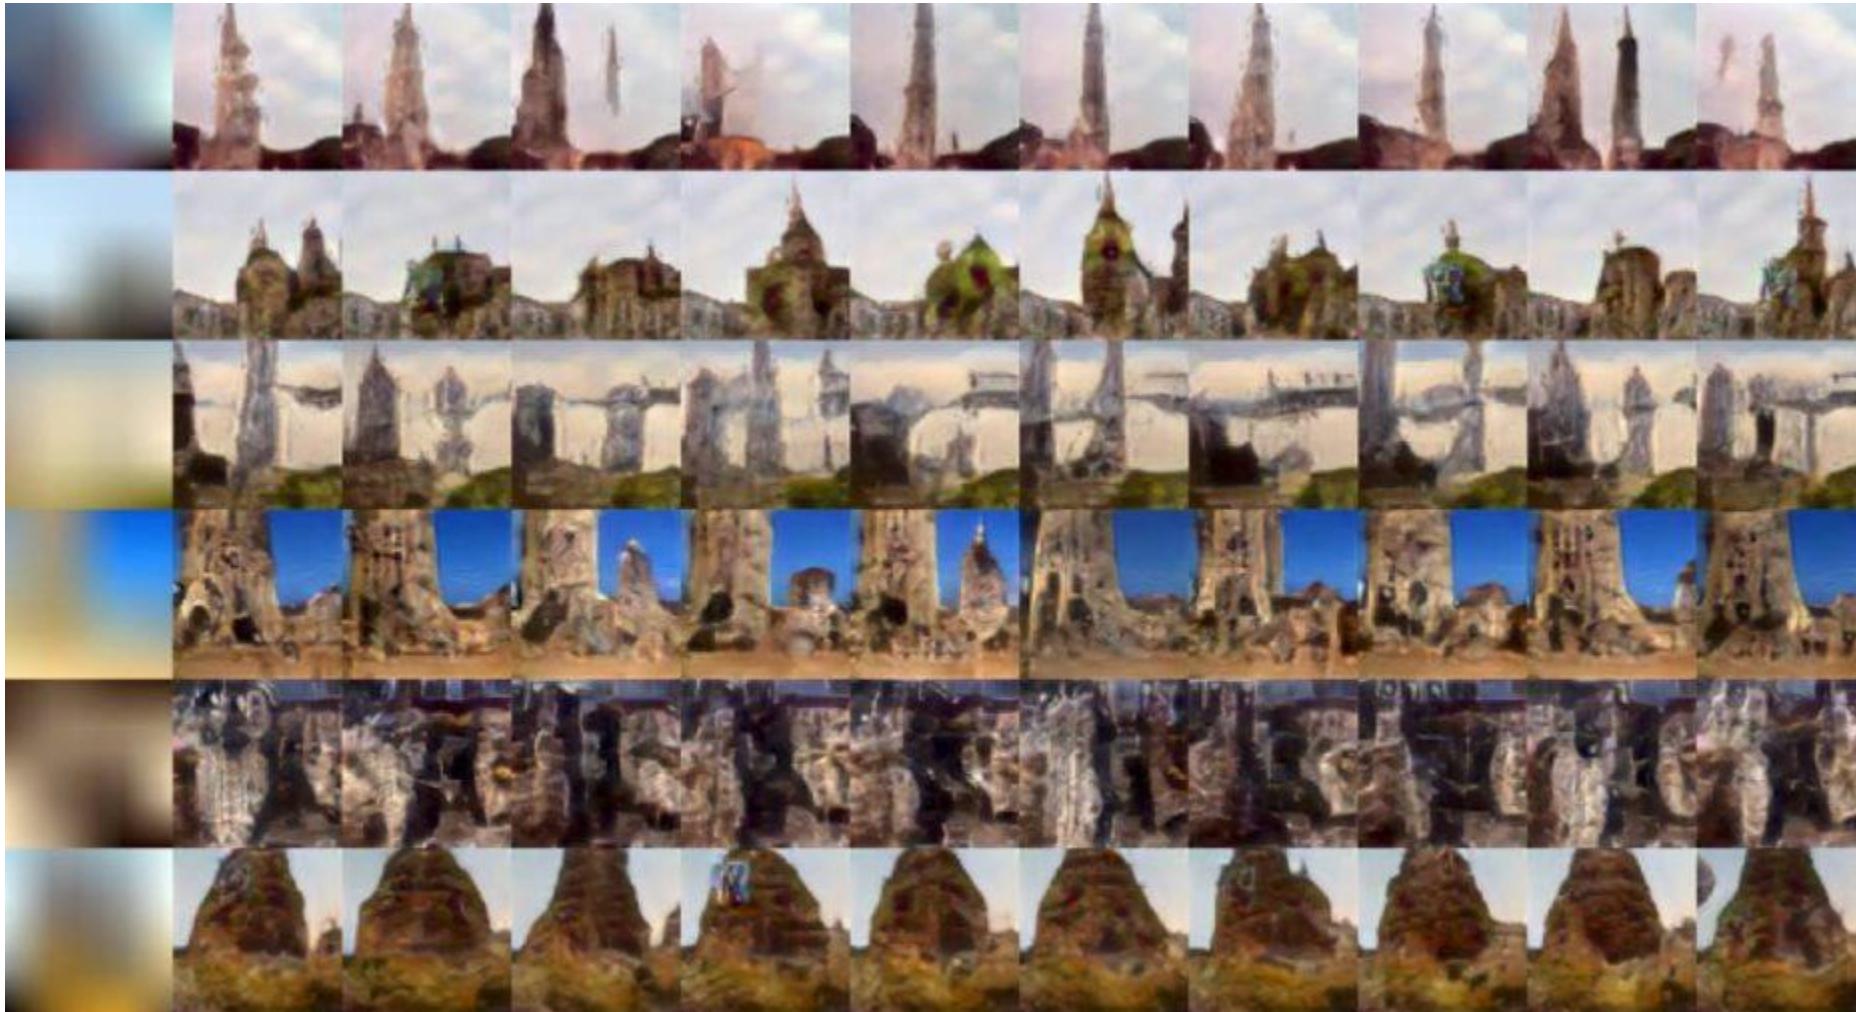
Example: Image captioning



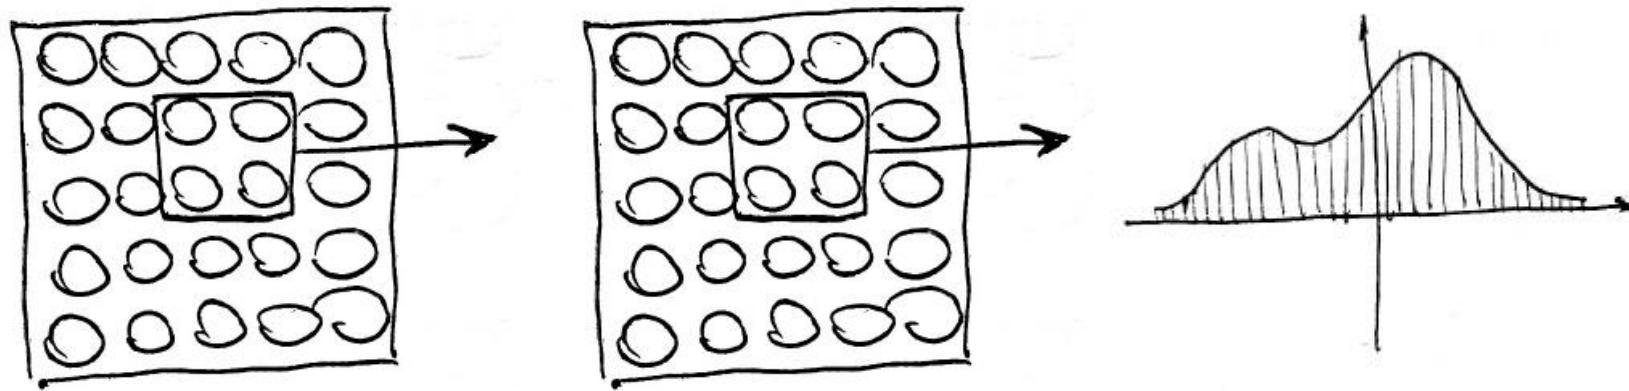
Example: Image captioning



Example: Image generation



Example: Image generation



Example: Sentiment analysis

“Thanks for a great party at the weekend, we really enjoyed it!”



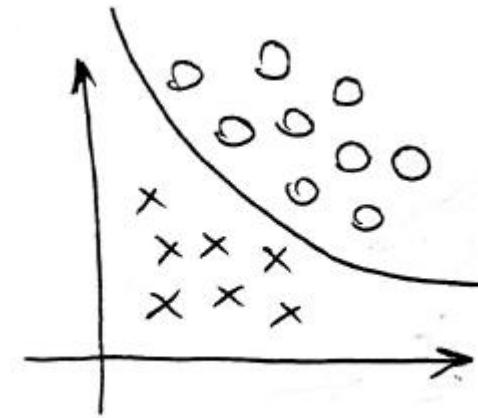
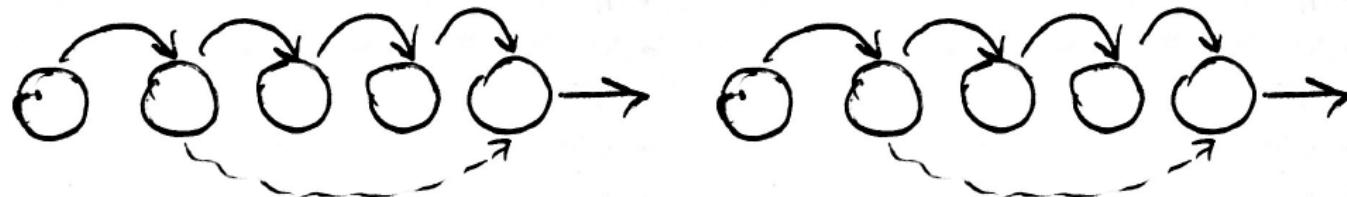
**sentiment: positive
score: 86%**

“I’m angry about the show, the acting was awful”

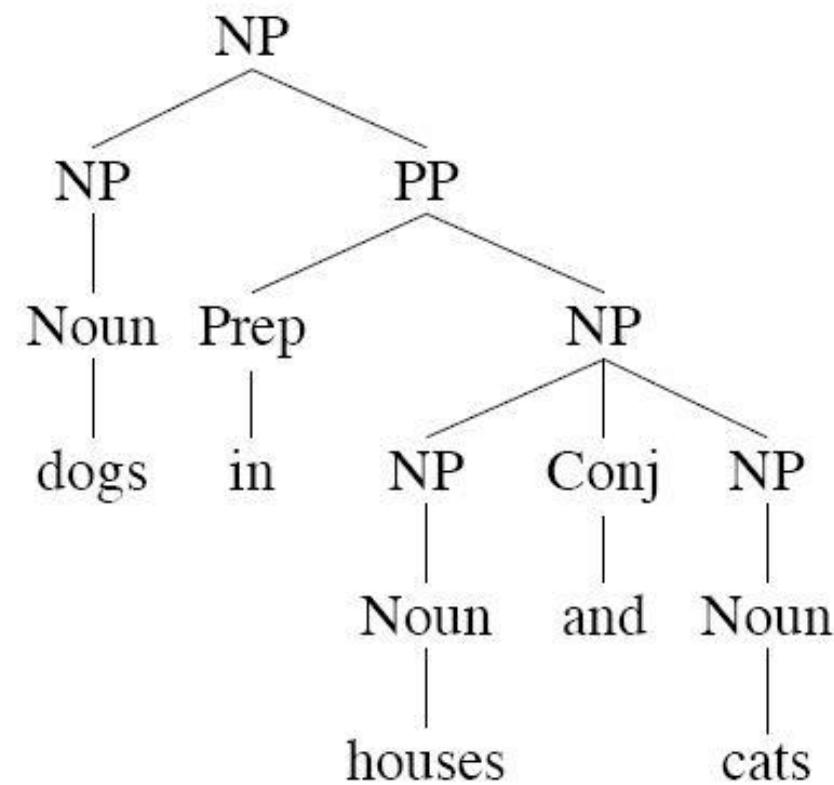
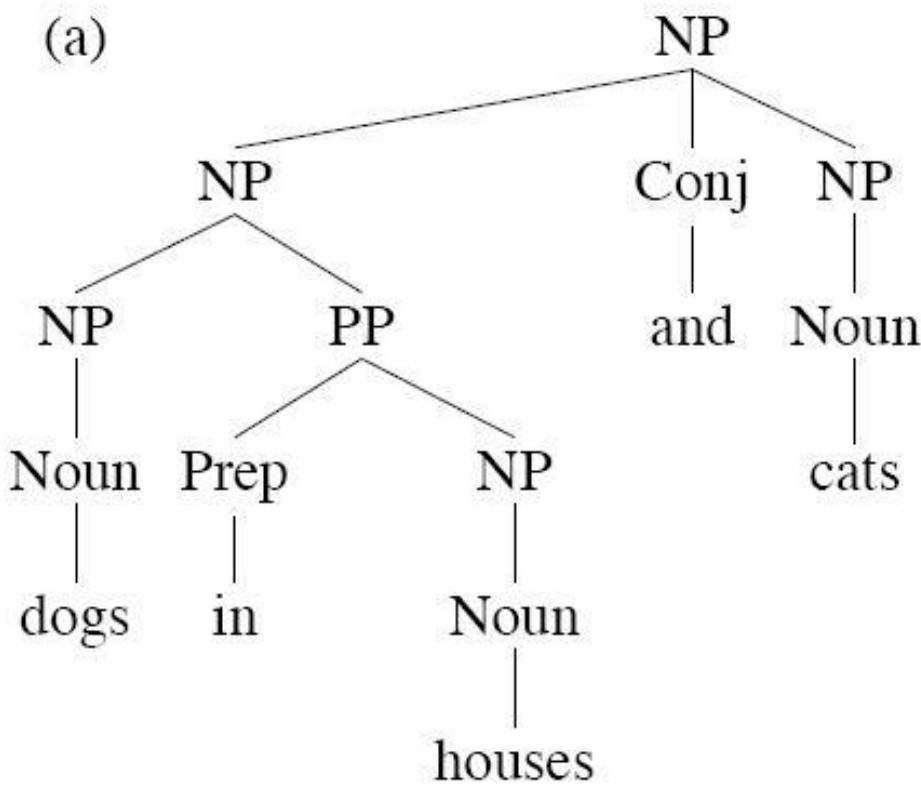


**sentiment: negative
score: -78%**

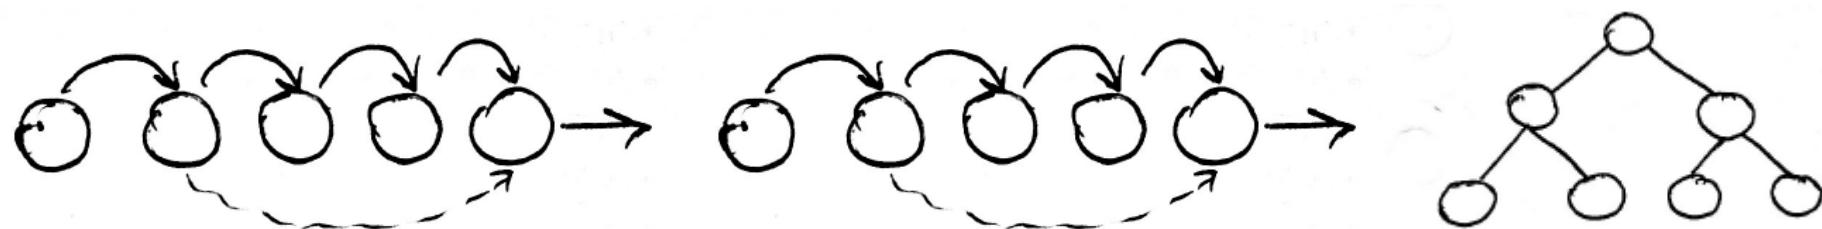
Example: Sentiment analysis



Example: Language understanding



Example: Language understanding



Example: Machine translation

Economic growth has slowed down in recent years .

Das Wirtschaftswachstum hat sich in den letzten Jahren verlangsamt .

Economic growth has slowed down in recent years .

La croissance économique s' est ralentie ces dernières années .

Example: Machine translation



Example: Neural conversational model

message Where do you live now?
response I live in Los Angeles.
message In which city do you live now?
response I live in Madrid.
message In which country do you live now?
response England, you?

message Where were you born?
response I was born in Canada.
message Where are you from?
response England, you?
message Where did you grow up?
response I grew up in Texas.

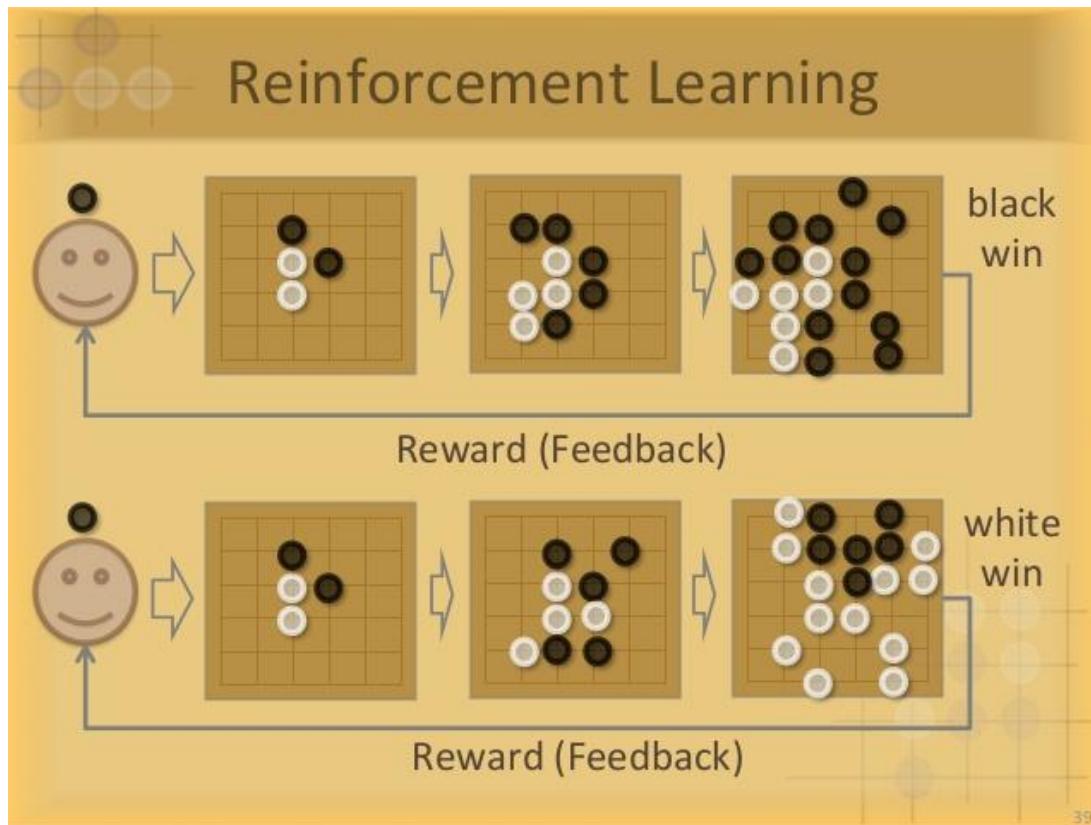
message How old are you?
response 16 and you?
message What's your age?
response 18.

message What is your major?
response I'm majoring in psychology
message What did you study in college?
response English lit.

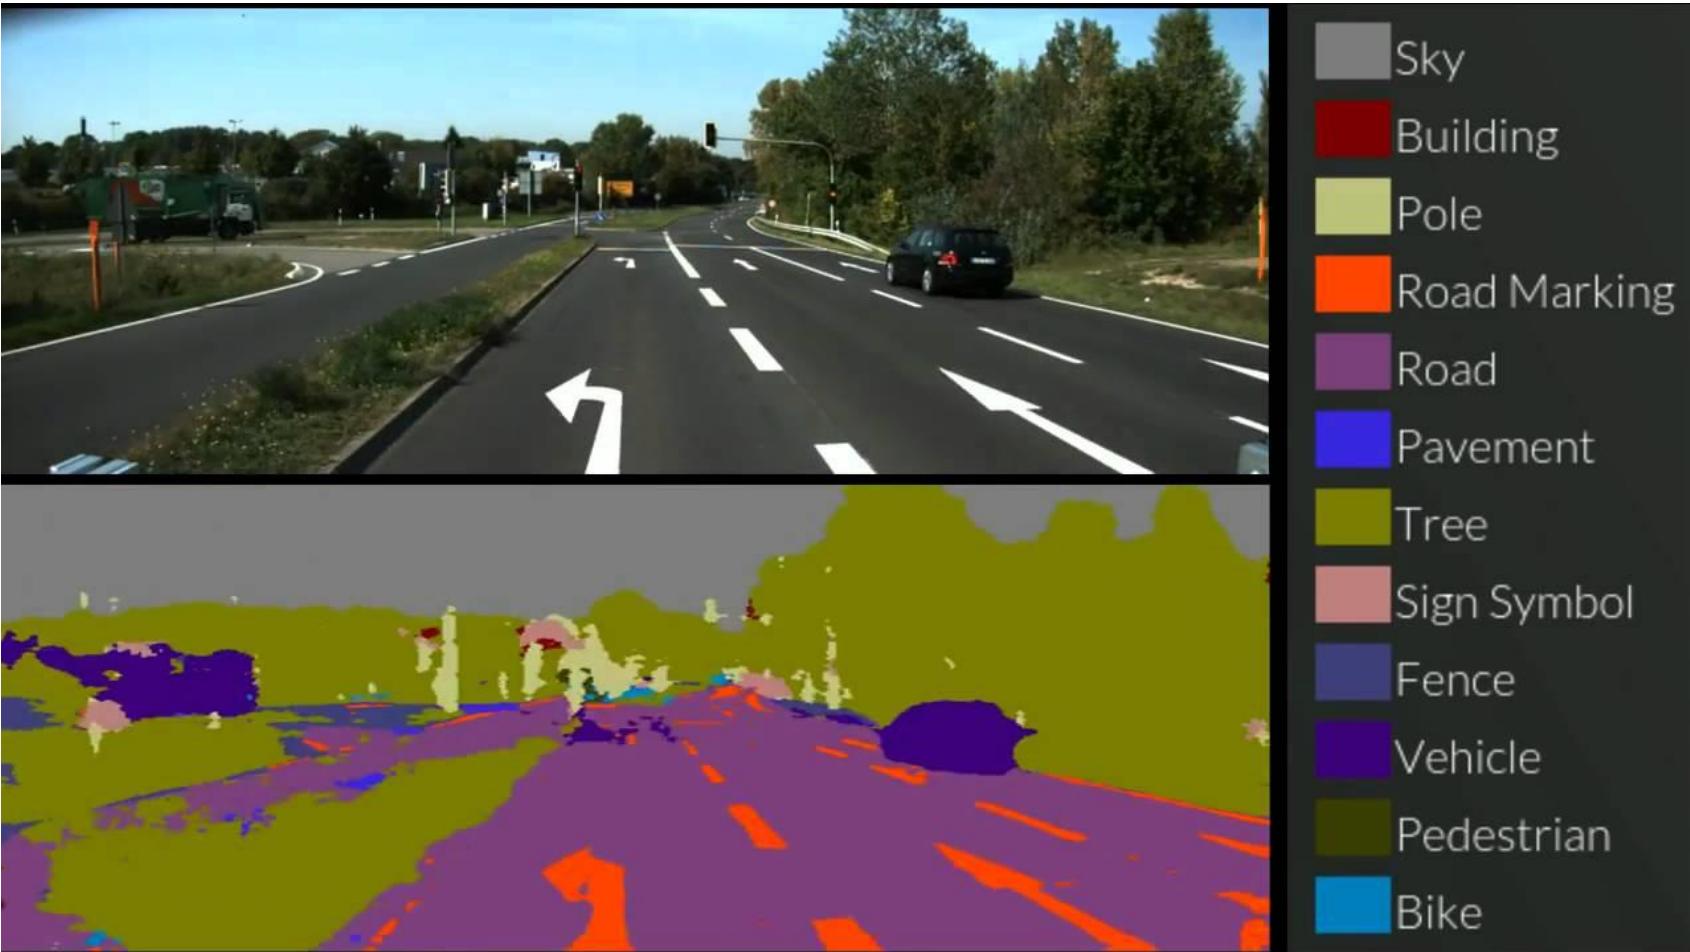
Example: Neural conversational model



Example: Playing games (reinforcement learning)



Example: Driving cars (reinforcement learning)



And a bit more...

- Video prediction (what will happen in the future!)
- Music generation
- Face recognition and search
- Emotion recognition
- Diagnostics in medicine
- Speech recognition and generation
- Artistic style transfer
- Autopilot car driving
- High frequency trading
- Neural photo editor
- Recommender systems
- Data noise reduction
- Visual question answering
- Robotics
- Chemistry, Physics...

Example: Object recognition

```
model = Sequential()

model.add(Convolution2D(nb_filters, kernel_size[0], kernel_size[1],
                      border_mode='valid',
                      input_shape=input_shape))
model.add(Activation('relu'))
model.add(Convolution2D(nb_filters, kernel_size[0], kernel_size[1]))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=pool_size))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer='adadelta',
              metrics=['accuracy'])

model.fit(X_train, Y_train, batch_size=batch_size, nb_epoch=nb_epoch,
          verbose=1, validation_data=(X_test, Y_test))
```

Example: Sentiment analysis

```
model = Sequential()
model.add(Embedding(max_features, 128, dropout=0.2))
model.add(LSTM(128, dropout_W=0.2, dropout_U=0.2)) # try
using a GRU instead, for fun
model.add(Dense(1))
model.add(Activation('sigmoid'))

# try using different optimizers and different optimizer configs
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

print('Train...')
model.fit(X_train, y_train, batch_size=batch_size,
nb_epoch=15,
validation_data=(X_test, y_test))
```

How it learns? (machine learning)

Program is said to learn from **experience E** (data set with pictures) with respect to some class of **tasks T** (logistic regression) and performance **measure P** (mean squared error) if performance improves (with gradient descent iterations)

Mitchell, 1997

How it learns? (deep learning)

Program is said to learn from learned representation R (**convolutional layers**) of experience E (**data set with pictures**) with respect to some class of tasks T (**logistic regression**) and performance measure P (**mean squared error**) if performance improves (**with gradient descent iterations**)

résumé

Deep learning networks are built from **blocks** – we combine and optimize these blocks to extract correct **representation** of complex data (images, text, speech) and give correct **output**

roadmap

Day 1 goals

- You **understand** basics of machine learning
- You **understand** difference between machine learning and deep learning (most engineers don't!)
- You **can** define machine learning problem with ERTP blocks

Day 2 goals

- You **understand** artificial neural networks structure
- You **understand** backpropagation algorithm (the most important thing)
- You **can** train your own neural network and continuously improve performance
- You **can** use Python **Keras** framework

Day 3 goals

- You **understand** how convolutional neural network (CNN) works
- You **understand** modern approaches to CNNs
- You **can** train your own CNN for different computer vision problems
- You **can** use **Caffe** framework

Day 4 goals

- You **understand** how recurrent neural network (RNN) works
- You **can** train your own RNN for text analysis
- You **can** use word2vec for text analysis
- You **can** train your own RNN for time series analysis
- You **can** use **TensorFlow** and **gensim** frameworks

Day 5 goals

- You **understand** modern research trends in deep learning
- You **understand** how to get a job in machine learning area
- You **can** perform the full deep learning pipeline after making a final project

Project proposals

- Titanic survivors prediction
 - House price prediction
- Fashion image classification
 - Food image classification
 - Arrhythmia ECG detection
 - Plants images classification
 - Adult content recognition
 - Stock price prediction
- Face identification
- Motivational quotes generation
 - Chat bots
 - Speech recognition
 - Machine translation
- Stealing ML models from API
 - Your ideas

thank you for attention!

rachnogstyle@gmail.com
facebook | rachnogstyle